

Lordphone Wen

New York, NY

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Large Language Model Systems, Machine Learning, Data Science for Software Engineering

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Large Language Models in Practice, Algorithms, Software Engineering

WORK EXPERIENCE

NewsBreak

May – August 2026

AI Engineer Intern

Mountain View, CA

- Shipped a personalized LLM news-alert pipeline (embed, LLM-score, LLM same-story dedupe) as a CronJob with a per-item decision log; hand-graded failures, ran prompt ablations and a 127-alert holdout, and added a categorical event gate, cutting filler 34% → 4% over 4,100+ shadow decisions.
- Held the alert pipeline behind a dry-run flag through 11 days of formal shadow reviews before live sends; per-user caps (2/day, 1/topic) held over 11,745 user-days with zero violations.
- Shipped streaming agentic search: RAG on turn one, then a bounded tool-calling agent (5 tools, ≤ 3 hops) over hybrid retrieval; server-resolved citations mean the model cannot author a URL.
- Took real logged queries that previously returned nothing from 0% to 100% P@1 with a deterministic entity resolver as the query rewriter (0.4 ms), measured by an eval harness run against the real catalog (28/28 recall).
- Diagnosed a 65% failure rate in another team's LLM pipeline, via Langfuse traces, as an unbounded reasoning loop hitting a 16,384-token ceiling; sized a 2,048-token cap against an observed p95 of 1,197.
- Built the Python/FastAPI and MongoDB backend from the first commit to 89 endpoints on EKS (Docker, Helm, Jenkins), with $\sim 1,450$ fully offline tests behind a ruff / pyright / pytest CI gate; moved nine third-party integrations server-side so the app ships zero API credentials.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Designed schemas and REST APIs for LangChain resume parsing, suggestions, and AI job matching.
- Led an 8-intern team through Agile sprints to ship a Next.js, FastAPI, and Supabase web app on GCP App Engine; led PR reviews and mentored interns on testing and Git workflows.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built an LLM mock-interview web app from scratch as the sole engineer (React, Next.js, Supabase); worked with recruiters so AI scoring matched real hiring criteria, and iterated scoring from user testing.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).
- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

SKILLS

Languages Python, SQL, TypeScript, JavaScript, Java, Go

LLM & ML RAG, hybrid retrieval, RRF, DeepSeek / DeepInfra, LangChain, Langfuse, PyTorch

Backend & Data FastAPI, Pydantic, pytest, MongoDB, PostgreSQL, MySQL, REST APIs

Cloud & DevOps AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, CI/CD

Certifications AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 20-frame (1s) video; processed 33 hours of Comma2k19 footage and ran streaming inference on local CUDA and GCP Vertex AI.