

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Large Language Model Systems, Machine Learning, Data Science for Software Engineering

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Large Language Models in Practice, Software Engineering, Algorithms

WORK EXPERIENCE

NewsBreak

May – August 2026

AI Engineer Intern

Mountain View, CA

- Built agentic LLM search: turn 1 is deterministic RAG over hybrid retrieval fused with RRF; follow-ups run a bounded tool-use agent (≤ 3 hops, 5 tools); server-resolved citations so the model cannot author a URL.
- Built an eval harness on real logged queries: entity resolution took P@1 from 0% to 100% at 0.4 ms after measurements ruled out embeddings; a manual grounding spot-check surfaced 23 answer problems.
- Shipped a chained LLM alert pipeline (embed $\rightarrow$ LLM-score $\rightarrow$ LLM same-story judge $\rightarrow$ push); hand-grading showed filler and real events both scored ~ 90 , so added a categorical gate; calibrated in production shadow on 4,100+ decisions and shipped live, cutting filler 34% $\rightarrow$ 4%.
- Diagnosed a 65% failure rate in another team's LLM pipeline via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways); sized a token cap that cuts 0% of successes.
- Built the FastAPI / MongoDB backend from its first commit to 89 endpoints on AWS EKS with $\sim 1,450$ offline tests behind a ruff / pyright / pytest CI gate and 344 PRs (340 merged); sole SwiftUI engineer for [SportsBreak](#) ($\sim 6,000$ installs), whose LLM digest reached 5,026 users.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Designed database schemas and REST APIs for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.
- Led an 8-intern team to take a React, Next.js, FastAPI, and Supabase web app from 0 $\rightarrow$ 1 on GCP App Engine; broke down user stories, led PR reviews, and mentored interns on Git workflows and testing.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0 $\rightarrow$ 1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned authentication, session state, and database interactions.
- Worked with recruiters so AI scoring categories matched real hiring, and iterated scoring and UI from user-testing sessions; demoed the product and won contracts with mid-sized companies.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

AI & LLMs

LangChain, RAG, hybrid retrieval, bge embeddings, Langfuse, DeepSeek / DeepInfra

Languages

Python, TypeScript, JavaScript, SQL, Swift, Go, Java, C, HTML, CSS

Frameworks & Data

FastAPI, Pydantic, React, Next.js, PostgreSQL, MongoDB, Supabase, Firebase, SwiftUI

Cloud & DevOps

AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, CI/CD

Practices

pytest, Unit Testing, REST APIs, Git, Agile/Scrum, Jira

Certifications

AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans:** Rating LLM coding plans bought from Z.AI, MiniMax, and Kimi; tested KV caching, long context, and needle-in-a-haystack to show which providers nerfed caching or heavily quantized models.