

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Machine Learning, Large Language Model Systems, Data Science for Software Engineering

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Large Language Models in Practice, Algorithms, Data Structure & Programming III

WORK EXPERIENCE

NewsBreak

May – August 2026

AI Engineer Intern

Mountain View, CA

- Shipped an agentic LLM search API: deterministic RAG on the first turn, a bounded agent with native tool calling on follow-ups (≤ 3 hops, 5 tools), hybrid retrieval fusing full-text and bge-embedding vector search with RRF, SSE streaming, and server-resolved citations the model cannot author.
- Designed and shipped a personalized LLM news-alert pipeline as a Kubernetes CronJob (ingest $\rightarrow$ embed $\rightarrow$ LLM-score $\rightarrow$ LLM same-story dedupe $\rightarrow$ push); evaluated it on 4,100+ production shadow decisions and hand-graded alerts, cutting filler 34% $\rightarrow$ 4%.
- Took real logged queries that previously returned nothing from 0% to 100% P@1 with a deterministic entity resolver as the query rewriter, after measuring embeddings against it on an eval harness run on the real catalog.
- Built the production Python FastAPI / MongoDB backend serving these features from its first commit to 89 REST endpoints on AWS EKS (Docker, Helm), with $\sim 1,450$ offline tests behind a ruff / pyright / pytest CI gate; wrote 34 contract-first API specs and $\sim 10,000$ lines of engineering docs.
- Diagnosed a 65% failure rate in another team's LLM pipeline, via Langfuse traces, as an unbounded reasoning loop hitting a 16,384-token ceiling; sized a 2,048-token cap against observed p95 of 1,197.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints and PR reviews to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; designed REST APIs for LangChain LLM resume parsing and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0 $\rightarrow$ 1 as the sole engineer (React, Next.js, Supabase, PostgreSQL); owned authentication, session state, and database interactions.
- Worked with recruiters so AI scoring categories matched real hiring and iterated scoring and UI from user-testing sessions; demoed the product and obtained contracts with mid-sized companies.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs; built an adaptive probabilistic safety certificate with PhD researchers that outperformed Adaptive MPC in simulation.

SKILLS

Languages Python, SQL, Java, JavaScript, TypeScript, Go, C

AI & ML RAG, bge embeddings, hybrid retrieval, RRF, LangChain, Langfuse, PyTorch, NumPy

Practices Git, CI/CD, REST APIs, Agile/Scrum, pytest, Unit Testing, TDD, ruff, pyright, Jira

Cloud & Data Docker, Kubernetes (EKS), Helm, AWS, GCP, FastAPI, MongoDB, PostgreSQL

Certifications AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 20-frame (1s) video; built a data pipeline over 33 hours of Comma2k19 and ran streaming inference on local CUDA and GCP Vertex AI.
- **ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.