

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Distributed Systems, Large Language Model Systems, Machine Learning

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Operating Systems, Large Language Models in Practice, Algorithms

WORK EXPERIENCE

NewsBreak

May – August 2026

AI Engineer Intern

Mountain View, CA

- Built agentic LLM search on one SSE endpoint: turn 1 is deterministic RAG over hybrid retrieval fused with RRF; follow-ups run a bounded tool-use agent (≤ 3 hops, 5 tools, 45s); server-resolved citations so the model cannot author a URL; entity resolution took real logged queries from 0% to 100% P@1.
- Shipped a personalized LLM alert pipeline (embed $\rightarrow$ LLM-score $\rightarrow$ LLM same-story judge $\rightarrow$ APNs) as a Kubernetes CronJob; calibrated it in production shadow and shipped live, cutting filler 34% $\rightarrow$ 4%.
- Debugged failures across the stack: a 65% failure rate in another team's async LLM workers, traced via Langfuse to an unbounded reasoning loop (p95 1,197 vs. 16,384-token runaways); a 5-day silent CronJob outage (ImagePullBackOff + Forbid concurrency); and a DNS $\rightarrow$ TLS $\rightarrow$ admission-policy chain blocking iOS clients.
- Built the async FastAPI / MongoDB (Motor) backend with AI coding agents from its first commit to 89 endpoints on AWS EKS, with $\sim 1,450$ offline tests behind a ruff / pyright / pytest CI gate; re-keyed an agent-written per-user cache to a shared one, $78\times$ smaller, as pod memory neared 512 MiB.
- Sole SwiftUI engineer for [SportsBreak](#) ($\sim 6,000$ installs, peak DAU 1,119) and authored 344 PRs (340 merged) across five repos; the app's LLM morning digest reached 5,026 users.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team to take a React, Next.js, FastAPI, and Supabase web app from 0 $\rightarrow$ 1 on GCP App Engine; broke down user stories, led PR reviews, and mentored interns on Git workflows and testing.
- Designed schemas and REST APIs for LangChain resume parsing, suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0 $\rightarrow$ 1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned authentication, session state, and database interactions.
- Worked with recruiters so AI scoring categories matched real hiring, and iterated scoring and UI from user-testing sessions; demoed the product and won contracts with mid-sized companies.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages	Python, TypeScript, JavaScript, SQL, Go, Java, Swift, C, HTML, CSS
Frameworks & Data	FastAPI, Pydantic, Motor / async MongoDB, React, Next.js, PostgreSQL, Supabase
AI & LLMs	RAG, hybrid retrieval, LangChain, Langfuse, bge embeddings, DeepSeek / DeepInfra
Cloud & DevOps	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, Grafana, Linux
Practices	pytest, ruff, pyright, Unit Testing, REST APIs, Git, Agile/Scrum
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av**: Built a PyTorch CNN-GRU that predicts steering and speed from 1s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI with streaming inference.