

EDUCATION

Carnegie Mellon University <i>M.S. in Software Engineering</i> GPA: 4.0 / 4.0 Selected Coursework: Large Language Model Systems, Distributed Systems, Machine Learning	Expected May 2027
University of Wisconsin–Madison <i>B.S. in Computer Science & B.S. in Computer Engineering</i> GPA: 3.72 / 4.0 Selected Coursework: Large Language Models in Practice, Software Engineering, Algorithms	May 2024

WORK EXPERIENCE

NewsBreak <i>AI Engineer Intern</i>	May – August 2026 <i>Mountain View, CA</i>
- Shipped agentic LLM search: turn 1 is deterministic RAG over hybrid retrieval fused with RRF; follow-ups run a bounded tool-use agent (≤ 3 hops, 5 tools); server-resolved citations so the model cannot author a URL. - Built an eval harness on real logged queries: entity resolution took P@1 from 0% to 100% at 0.4 ms after measurements ruled out embeddings; a manual grounding spot-check surfaced 23 answer problems. - Shipped a personalized LLM alert pipeline (embed $\rightarrow$ LLM-score $\rightarrow$ LLM same-story judge $\rightarrow$ push); hand-grading showed filler and real events both scored ~ 90, so added a categorical gate; calibrated in production shadow on 4,100+ decisions and shipped live, cutting filler 34% $\rightarrow$ 4%. - Diagnosed a 65% failure rate in another team’s LLM pipeline via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways); sized a token cap that cuts 0% of successes. - Built the FastAPI / MongoDB backend from its first commit to 89 endpoints on AWS EKS with $\sim 1,450$ offline tests behind a ruff / pyright / pytest CI gate and 344 PRs (340 merged); sole SwiftUI engineer for SportsBreak ($\sim 6,000$ installs), whose LLM digest reached 5,026 users.	
Haddee Education <i>Engineering Lead</i>	June – September 2025 <i>San Diego, CA</i>
- Designed database schemas and REST APIs for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching. - Led an 8-intern team to take a React, Next.js, FastAPI, and Supabase web app from 0$\rightarrow$1 on GCP App Engine; broke down user stories, led PR reviews, and mentored interns on Git workflows and testing.	
ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch <i>Full Stack Developer</i>	September – December 2023 <i>Beijing, China</i>
- Built a VC-backed startup’s LLM mock-interview web app from 0$\rightarrow$1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned authentication, session state, and database interactions. - Worked with recruiters so AI scoring matched real hiring; demoed the product to clients, ran fast iterations on their feedback, and won contracts with mid-sized companies.	

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University <i>Graduate Researcher</i> Advisor: Prof. Yorie Nakahira	February – June 2026
Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).	

SKILLS

AI & LLMs	RAG, hybrid retrieval, LangChain, Langfuse, bge embeddings, DeepSeek / DeepInfra
Languages	Python, TypeScript, Go, JavaScript, SQL, Swift, Java, C, HTML, CSS
Frameworks & Data	React, Next.js, FastAPI, Pydantic, PostgreSQL, MongoDB, Supabase, Firebase, SwiftUI
Cloud & DevOps	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, CI/CD
Practices	pytest, Unit Testing, REST APIs, Git, Agile/Scrum, Jira
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

codingplans: Rating LLM coding plans bought from Z.AI, MiniMax, and Kimi; tested KV caching, long context, and needle-in-a-haystack to show which providers nerfed caching or heavily quantized models.
--