

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Machine Learning, Large Language Model Systems, Data Science for Software Engineering

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Large Language Models in Practice, Algorithms, Data Structure & Programming III

WORK EXPERIENCE

NewsBreak

May – August 2026

AI Engineer Intern

Mountain View, CA

- Designed a personalized breaking-news classifier: bge embeddings, LLM scoring, a per-topic 7-day percentile gate, and an LLM same-story judge (~1,200 LLM calls/day).
- Calibrated it in production shadow on 4,100+ decision rows, then shipped it live: filler 34% → 4% by count (29% → 3% device-weighted), with the 2/day cap held over 11,745 user-days.
- Hand-graded alerts and found filler and real events both scored 89–90, so added a categorical gate; rejected a rule that caught 91% of filler but killed the day's two highest-reach alerts.
- Shipped agentic search with hybrid retrieval (Mongo text + bge vectors, RRF) and a bounded tool-calling agent; took logged queries from 0% to 100% P@1 with 0.4ms entity resolution.
- Traced another team's 65% LLM failure rate via Langfuse traces to an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways); a binomial model predicted 49–66% failure vs. 65% observed.
- Built the FastAPI / MongoDB backend from its first commit to 89 endpoints on AWS EKS, with ~1,450 offline tests; cut a per-user cache 78× (13.2MB → 10.9MB shared) to fit a 512MiB pod.
- Sole SwiftUI engineer on [SportsBreak](#) (1.4K App Store ratings, 1,119 peak daily active users); shipped a personalized LLM morning digest to 5,026 users (~81% of the base).
- On [DramaBreak](#), ran an A/B test of menu options, measuring click-through rate in Amplitude.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; designed REST APIs for LangChain resume parsing and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer; iterated AI scoring from user-testing sessions and with recruiters so scoring categories matched real hiring.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).
- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

SKILLS

Languages Python, Java, SQL, Go, C, TypeScript, JavaScript, Swift, Bash

ML & LLM PyTorch, CUDA, NumPy, RAG, hybrid retrieval, RRF, bge embeddings, Bayesian estimation

AI Tooling Langfuse, LangChain, LLM tool calling / agents, Claude Code / AI coding tools

Data & Infra MongoDB, PostgreSQL, FastAPI, Kubernetes (EKS), Docker, AWS, GCP Vertex AI, CI/CD

PROJECTS

- my-av:** Developed a PyTorch CNN-GRU for real-time steering-angle and speed prediction from 20-frame (1s) video; built a data pipeline over 33 hours of Comma2k19 with temporal windowing, trained on local CUDA and GCP Vertex AI, and implemented streaming inference.