

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Distributed Systems, Foundations of Software Engineering, Large Language Model Systems

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Algorithms, Data Structure & Programming III, Operating Systems, Software Engineering

WORK EXPERIENCE

NewsBreak

May – August 2026

Backend Engineer Intern

Mountain View, CA

- Built SportsBreak's FastAPI / async MongoDB backend from its first commit to 89 REST endpoints and 26 collections (~83% of backend commits), with ~1,450 offline unit tests behind a ruff / pyright / pytest CI gate.
- Stood up production on AWS EKS (Docker, Helm, Jenkins) on day 3 with HPA autoscaling and 7 CronJobs, and set up Prometheus metrics and Grafana dashboards for the backend.
- Root-caused a 5-day silent CronJob outage (a non-terminal ImagePullBackOff held the Job active, so Forbid concurrency skipped every later run) and fixed it with paired job and scheduler deadlines.
- Replaced a 25–30 request client fan-out with one backend call, and redesigned a per-user cache that could never hit (13.2 MB per user) into a shared 10.9 MB catalog, 78× smaller, to fit a 512 MiB pod.
- Shipped Stripe web checkout for DramaBreak with a signature-verified webhook as the only trusted fulfillment path, made idempotent with a 60 s stale-pending reclaim.
- Designed an agentic search API over SSE: deterministic RAG on turn 1, then a bounded tool-calling agent (≤ 3 hops, 5 tools); hybrid retrieval (Mongo text + vector, RRF) degrades to surviving legs.
- Designed an async ingest $\rightarrow$ embed $\rightarrow$ LLM-score $\rightarrow$ dedupe $\rightarrow$ push pipeline (~1,200 LLM calls/day); calibrated on 4,100+ shadow decisions, it cut filler 34% $\rightarrow$ 4% and held its 2/day cap over 11,745 user-days.
- Traced another team's 65% LLM failure rate via Langfuse telemetry to an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways) and a Helm env shadowing the code default.
- Sole SwiftUI engineer for [SportsBreak](#) (1.4K App Store ratings, 1,119 peak daily active users); wrote 34 contract-first API specs and ~10,000 lines of engineering docs, and onboarded the second engineer.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0 $\rightarrow$ 1 on GCP App Engine; designed the database schemas and REST APIs for its LangChain LLM features.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing protocols, and scalable code practices.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0 $\rightarrow$ 1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned auth, session state, and database interactions.
- Worked with recruiters to match AI scoring to real hiring, then demoed and iterated with clients until three 200–500-person companies signed contracts.

SKILLS

Languages	Python, Java, JavaScript, TypeScript, Go, C, SQL, Swift, Bash
Backend & Data	REST APIs, SSE, FastAPI, asyncio, MongoDB, PostgreSQL, MySQL, Redis Streams
Cloud & Observability	AWS, Kubernetes (EKS), Docker, Helm, Jenkins, Prometheus, Grafana, Langfuse
AI & Frameworks	LLM tool calling / agents, RAG, hybrid retrieval, LangChain, React, Next.js, SwiftUI
Practices	Unit Testing, pytest, TDD, CI/CD, Git, Agile/Scrum, Claude Code / AI coding tools
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av**: Built a PyTorch CNN-GRU that predicts steering and speed from 1s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI.