

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Distributed Systems, Foundations of Software Engineering, Large Language Model Systems

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Algorithms, Data Structure & Programming III, Operating Systems, Software Engineering

WORK EXPERIENCE

NewsBreak

May – August 2026

Backend Engineer Intern

Mountain View, CA

- Built SportsBreak's FastAPI / async MongoDB backend from its first commit to 89 REST endpoints and 26 collections, with ~1,450 offline unit tests behind a ruff / pyright / pytest CI gate; authored 344 PRs.
- Stood up production Kubernetes (AWS EKS, Helm, Jenkins) on day 3 with 7 CronJobs; fixed a 5-day silent CronJob outage (ImagePullBackOff + Forbid concurrency) with paired job and scheduler deadlines.
- Redesigned a cache that could not hit (13.2MB keyed per user) into a shared 10.9MB catalog cache, 78× smaller and sized for a 512 MiB pod; replaced a 25–30 request client fan-out with one backend call.
- Designed an agentic search API: one SSE endpoint, deterministic RAG on turn 1 and a bounded tool-calling agent (≤ 3 hops, 5 tools) after; hybrid retrieval (Mongo text + vector, RRF) degrades to surviving legs.
- Designed an async ingest $\rightarrow$ embed $\rightarrow$ LLM-score $\rightarrow$ dedupe $\rightarrow$ push pipeline (~1,200 LLM calls/day); calibrated on 4,100+ shadow decisions, it cut filler 34% $\rightarrow$ 4% and held its 2/day cap over 11,745 user-days.
- Traced another team's 65% LLM failure rate via Langfuse telemetry to an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways); a binomial model predicted 49–66% failure vs. 65% observed.
- Sole SwiftUI engineer; shipped [SportsBreak](#) to the App Store (~6,000 installs, 1,119 peak daily active users); wrote 34 contract-first API specs and ~10,000 lines of engineering docs, and onboarded the second engineer.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0 $\rightarrow$ 1 on GCP App Engine; designed the schemas and REST APIs for LangChain resume parsing.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing protocols, and scalable code practices.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0 $\rightarrow$ 1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned auth, session state, and database interactions.
- Demoed the product and iterated with clients until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

SKILLS

Languages	Python, Java, JavaScript, TypeScript, Go, C, SQL, Swift, Bash
Backend & Data	REST APIs, SSE, FastAPI, asyncio, MongoDB, PostgreSQL, MySQL, Redis Streams
Cloud & Infra	Docker, Kubernetes (EKS), Helm, AWS, GCP, Jenkins, GitHub Actions, CI/CD, Langfuse
AI & Frameworks	LLM tool calling / agents, RAG, hybrid retrieval, LangChain, React, Next.js, SwiftUI
Practices	Unit Testing, pytest, TDD, Git, Agile/Scrum, Claude Code / AI coding tools

PROJECTS

- my-av:** Built a PyTorch CNN-GRU that predicts steering and speed from 1s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI.