

Lordphone Wen

New York, NY

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Large Language Model Systems, Distributed Systems, Foundations of Software Engineering

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Software Engineering, Large Language Models in Practice, Operating Systems, Algorithms

WORK EXPERIENCE

NewsBreak

May – August 2026

Software Engineering Intern

Mountain View, CA

- Built the FastAPI / MongoDB backend with AI coding tools from its first commit to 89 endpoints and ~1,450 fully-offline tests (0.17 s); added a ruff / pyright / pytest CI gate after two tests sat red on `main`.
- Stood up production AWS EKS (Helm, Jenkins) on day 3 with 7 CronJobs and a two-stage amd64 image; diagnosed a DNS → TLS → admission-policy chain blocking iOS clients; caught a stale ArgoCD app.
- Diagnosed a 5-day silent CronJob outage (non-terminal ImagePullBackOff + Forbid concurrency) and fixed it with paired job and scheduler deadlines; traced a 65% LLM failure rate via Langfuse.
- Automated a personalized LLM alert pipeline in one CronJob (embed → LLM-score → LLM same-story judge → push); calibrated it in production shadow and shipped it, cutting filler 34% → 4%.
- Shipped agentic LLM search: deterministic RAG over hybrid retrieval (keyword, full-text, bge embeddings, RRF), a bounded tool-use agent, and server-resolved citations; entity resolution hit 100% P@1.
- Sole SwiftUI engineer for [SportsBreak](#) (~6,000 installs, peak DAU 1,119) and primary backend author; 344 PRs (340 merged); wrote 34 contract-first API specs and ~10,000 lines of docs, and onboarded the next engineer.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; designed the schemas and REST APIs for its LangChain LLM features.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing protocols, and scalable code practices.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned authentication, session state, and database interactions.
- Worked with recruiters so AI scoring matched real hiring; demoed the product, iterated fast on client feedback, and won contracts with three 200–500-person companies.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages	Python, TypeScript, JavaScript, SQL, Go, Java, Swift, C, HTML, CSS
Cloud & DevOps	AWS, Terraform, Docker, Kubernetes (EKS), Helm, Jenkins, GitHub Actions, CI/CD
Frameworks & Data	React, Next.js, Express, FastAPI, PostgreSQL, MongoDB, Redis Streams, Supabase
AI & LLMs	Claude Code / AI coding tools, LLM tool calling / agents, RAG, Langfuse, LangChain
Practices	pytest, Unit Testing, ruff, pyright, REST APIs, Git, Agile/Scrum
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans**: Rating LLM coding plans bought from Z.AI, MiniMax, and Kimi; tested KV caching, long context, and needle-in-a-haystack to show which providers nerfed caching or heavily quantized models.