

EDUCATION

Carnegie Mellon University <i>M.S. in Software Engineering GPA: 4.0 / 4.0</i> Selected Coursework: Large Language Model Systems, Distributed Systems, Foundations of Software Engineering	Expected May 2027
University of Wisconsin–Madison <i>B.S. in Computer Science & B.S. in Computer Engineering GPA: 3.72 / 4.0</i> Selected Coursework: Large Language Models in Practice, Software Engineering, Algorithms, Operating Systems	May 2024

WORK EXPERIENCE

NewsBreak <i>Software Engineering Intern</i>	May – August 2026 <i>Mountain View, CA</i>
- Shipped agentic LLM search on one SSE endpoint: turn 1 is deterministic RAG over hybrid retrieval (keyword, full-text, and bge vectors fused with RRF); follow-ups run a bounded tool-use agent (≤ 3 hops, 5 tools, 45 s); server-resolved citations so the model cannot author a URL. - Built an eval harness on real logged queries: deterministic entity resolution took P@1 from 0% to 100% at 0.4 ms after measurements ruled out embeddings; a manual grounding spot-check surfaced 23 answer problems. - Shipped a personalized LLM alert pipeline (embed $\rightarrow$ LLM-score $\rightarrow$ LLM same-story judge $\rightarrow$ push); added a categorical gate after hand-grading, calibrated in production shadow, and shipped live, cutting filler 34% $\rightarrow$ 4%. - Built the FastAPI / MongoDB backend with AI coding tools from its first commit to 89 endpoints on AWS EKS, with $\sim 1,450$ offline tests behind a ruff / pyright / pytest CI gate; re-keyed an agent-written per-user cache to a shared one, $78\times$ smaller, as pod memory neared 512 MiB. - Moved nine third-party integrations server-side so the iOS binary ships zero API credentials; reopened two closed security tickets with line-level evidence that the fixes were absent. - Sole SwiftUI engineer for SportsBreak ($\sim 6,000$ installs, peak DAU 1,119) and primary backend author; 344 PRs (340 merged) across five repos; its LLM morning digest reached 5,026 users.	
Haddee Education <i>Engineering Lead</i>	June – September 2025 <i>San Diego, CA</i>
Led an 8-intern team to ship a React, Next.js, FastAPI, and Supabase web app from 0$\rightarrow$1 on GCP App Engine; designed the schemas and REST APIs for its LangChain LLM features and led PR reviews.	
ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch <i>Full Stack Developer</i>	September – December 2023 <i>Beijing, China</i>
- Built a VC-backed startup’s LLM mock-interview web app from 0$\rightarrow$1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned authentication, session state, and database interactions. - Worked with recruiters so AI scoring matched real hiring; demoed the product, iterated fast on client feedback, and won contracts with three 200–500-person companies.	

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University <i>Graduate Researcher Advisor: Prof. Yorie Nakahira</i>	February – June 2026
Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).	

SKILLS

Languages	Python, TypeScript, JavaScript, Go, Java, SQL, Swift, C, HTML, CSS
AI & LLMs	RAG, hybrid retrieval, Claude Code / AI coding tools, LLM tool calling / agents
Frameworks & Data	React, Next.js, FastAPI, Pydantic, PostgreSQL, MongoDB, Supabase, Firebase, SwiftUI
Cloud & DevOps	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, CI/CD
Practices	pytest, Unit Testing, REST APIs, Server-Sent Events (SSE), Git, Agile/Scrum
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

codingplans: Rating LLM coding plans bought from Z.AI, MiniMax, and Kimi; tested KV caching, long context, and needle-in-a-haystack to show which providers nerfed caching or heavily quantized models.	
--	--