

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Large Language Model Systems, Distributed Systems, Machine Learning, Data Science for Software Engineering

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Large Language Models in Practice, Software Engineering, Algorithms, Operating Systems

WORK EXPERIENCE

NewsBreak (Venture Team)

May – August 2026

AI Engineer Intern

Mountain View, CA

- Designed and shipped a personalized LLM alert pipeline as a Kubernetes CronJob (embed → LLM-score → select → LLM same-story dedupe → push; ~1,200 LLM calls/day); calibrated it on 4,100+ logged decisions in production shadow before going live; filler 34% → 4%.
- Shipped agentic search on one SSE endpoint: turn 1 is deterministic RAG over hybrid retrieval (RRF); follow-ups run a bounded tool-calling agent (≤3 hops, 5 tools, 45s) with server-resolved citations the model cannot author; entity resolution hit 100% P@1 on real logged queries.
- Owned API integrations: migrated nine third-party integrations (SportsDataIO, YouTube, SerpApi, DeepInfra LLM + embeddings, APNs) server-side so the iOS binary ships zero API credentials; shipped Stripe web checkout behind a signature-verified, idempotent webhook.
- Built the FastAPI / MongoDB backend from the first commit to 89 endpoints on AWS EKS, with ~1,450 fully-offline tests behind a ruff / pyright / pytest CI gate; stood up production EKS on day 3.
- Built a live manual segment-broadcast tool used during the World Cup, with alerts landing a minute before FOX broadcast the goal; introduced the team's Jira backlog unprompted and wrote 34 contract-first API specs.
- Diagnosed a 65% failure rate in another team's LLM pipeline via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways) and sized a 2,048-token cap against observed successes.
- Sole SwiftUI engineer for [SportsBreak](#) (~6,000 installs) and primary backend author; 344 PRs, 340 merged.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Designed database schemas and REST APIs for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.
- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; broke down user stories and led PR reviews.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase), built from 0→1; worked with recruiters so AI scoring categories matched real hiring, and owned auth, session state, and database interactions.
- Demoed the product and iterated with clients until three 200–500-person companies signed contracts.

SKILLS

AI & Agents	LLM tool calling / agents, RAG, LangChain, MCP, Langfuse, Claude Code / AI coding tools
Languages	Python, TypeScript, JavaScript, SQL, Go, Java, Swift
Backend & APIs	FastAPI, REST APIs, Server-Sent Events (SSE), Stripe, Supabase, PostgreSQL, MongoDB
Frontend	React, Next.js, SwiftUI
Infra & CI	AWS, GCP, Kubernetes (EKS), Docker, GitHub Actions, CI/CD, Git

PROJECTS

- codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and heavy quantization.
- my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 1s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI, with streaming inference.