

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Distributed Systems, Large Language Model Systems, Machine Learning

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Operating Systems, Computer Architecture, Mobile Computing Laboratory

WORK EXPERIENCE

NewsBreak

May – August 2026

Software Engineering Intern

Mountain View, CA

- Shipped a multi-step LLM alert pipeline as a Kubernetes CronJob (embed → LLM-score → select → LLM dedupe → push; ~1,200 calls/day); shadowed 4,100+ decisions in production before going live; filler 34% → 4%.
- Shipped agentic search on one SSE endpoint: turn 1 is deterministic RAG over hybrid retrieval (RRF); follow-ups run a bounded tool-use agent (≤ 3 hops, 5 tools, 45 s); a failed retrieval leg degrades to the survivors, and server-resolved citations mean the model cannot author a URL.
- Gated AI-written code with ~1,450 offline tests (0.17 s) behind a ruff / pyright / pytest CI gate, using a Motor fake that reproduces a prod failure that silently killed a search leg behind HTTP 200.
- Redesigned a cache that could not hit after pod memory neared its 512 MiB limit: 13.2 MB per user → 10.9 MB shared (78×), serving one backend call in place of a 25–30 request client fan-out.
- Built server-driven UI over 27 widget types so a new catalog surface ships without an App Store release and unknown widget kinds drop instead of failing a page; shipped a 72-event typed analytics catalog with a server-side allowlist.
- Diagnosed a 65% failure rate in another team's LLM pipeline via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways); found iOS silently drops a Live Activity start push while APNs returns 200.
- Sole SwiftUI engineer for [SportsBreak](#) (~6,000 installs) and primary backend author; stood up production AWS EKS on day 3 and fixed a 5-day silent CronJob outage.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team to ship a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; designed schemas and REST APIs for LangChain resume parsing and AI job matching.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing protocols, and scalable code practices.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase), built from 0→1; iterated AI scoring and UI with clients until three 200–500-person companies signed.

SKILLS

Systems & Infra	Kubernetes (EKS), CronJobs, Docker, Helm, Redis Streams, AWS, NVIDIA Jetson
AI & Agents	LLM tool calling / agents, Claude Code / AI coding tools, MCP, RAG, Langfuse
Languages	Python, C, Go, Swift, TypeScript, JavaScript, Java, SQL
Frameworks & Data	FastAPI, asyncio, Server-Sent Events (SSE), SwiftUI, React, Next.js, MongoDB
Testing & CI	pytest, ruff, pyright, GitHub Actions, Jenkins, CI/CD, Git

PROJECTS

- Wisconsin Autonomous:** Integrated wheel-speed encoders with voltage-to-frequency converters on the competition rover and streamed real-time data to the onboard NVIDIA Jetson.
- my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 1 s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI, with streaming inference.
- codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and heavy quantization.