

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Large Language Model Systems, Machine Learning, Data Science for Software Engineering, Distributed Systems

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Large Language Models in Practice, Software Engineering, Algorithms

WORK EXPERIENCE

NewsBreak

May – August 2026

AI Engineer Intern

Mountain View, CA

- Designed and shipped a personalized LLM alert workflow as a Kubernetes CronJob (embed → LLM-classify and score → per-topic percentile select → LLM same-story dedupe → push; ~1,200 LLM calls/day), with a decision log recording why each item fired or did not.
- Evaluated it over 11 days of production shadow on 4,100+ logged decisions: hand-graded alerts showed filler and real events both scored 89–90, so added a categorical *happened* / *speculative* / *no_event* gate, checked it on a 127-alert holdout, and ran prompt ablations; rejected a rule that caught 91% of filler but blocked the day's two highest-reach real alerts; shipped live with filler 34% → 4%.
- Shipped RAG and agentic search on one SSE endpoint: turn 1 is deterministic RAG over hybrid retrieval (substring + text + vector, fused by RRF); follow-ups run a bounded tool-calling agent limited to 5 whitelisted tools (≤3 hops, 45 s), with server-resolved citations the model cannot author.
- Built an eval harness on real logged queries against the committed catalog; deterministic entity resolution took P@1 from 0% to 100% (28/28 recall) at 0.4 ms, after measurements ruled out embeddings.
- Diagnosed a 65% failure rate in another team's LLM pipeline via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways) and sized a 2,048-token cap against observed successes.
- Built the FastAPI / MongoDB backend from the first commit to 89 endpoints on AWS EKS, with ~1,450 fully-offline tests behind a ruff / pyright / pytest CI gate; moved nine third-party integrations (LLM, embeddings, data APIs) server-side so the client ships zero API credentials.
- Sole SwiftUI engineer for [SportsBreak](#) (~6,000 installs); shipped a personalized LLM morning digest used by 5,026 users (~81% of the base), tracked through a 72-event Amplitude analytics catalog.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Designed database schemas and REST APIs for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.
- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Sole engineer on a VC-backed startup's LLM mock-interview app (React, Next.js, Supabase); shaped AI scoring with recruiters and iterated from user testing until three 200–500-person companies signed.

SKILLS

AI & ML

LLM tool calling / agents, RAG, hybrid retrieval (RRF), LangChain, Langfuse, PyTorch

Languages

Python, TypeScript, SQL, JavaScript, Go, Swift

Full Stack & Data

FastAPI, REST APIs, SSE, React, Next.js, PostgreSQL, MongoDB, Supabase

Cloud & Observability

AWS, GCP, Kubernetes (EKS), Docker, GitHub Actions, CI/CD, Amplitude

Practices

pytest, Unit Testing, Git, Jira, Agile/Scrum, Claude Code / AI coding tools

PROJECTS

- codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and heavy quantization.