

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Distributed Systems, Large Language Model Systems, Machine Learning

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Operating Systems, Computer Architecture, Large Language Models in Practice

WORK EXPERIENCE

NewsBreak

May – August 2026

Engineering Intern

Mountain View, CA

- Built a personalized LLM scoring pipeline as a Kubernetes CronJob (bge embeddings → LLM scoring → per-topic 7-day percentile gate → LLM dedupe → push, ~1,200 LLM calls/day) with a per-item decision log.
- Ramped it through 11 days of production shadow (4,100+ logged decisions) behind a dry-run flag, then shipped it live: filler 34% → 4% by count, with the 2/day cap held over 11,745 user-days.
- Root-caused a 65% failure rate in another team's LLM workers via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways) and a Helm env shadowing the code default.
- Shipped agentic search with hybrid retrieval (Mongo text + bge vectors fused by RRF) and a bounded tool-calling agent over SSE, taking logged queries from 0% to 100% P@1 with 0.4 ms entity resolution.
- Stood up a production service on AWS EKS (Docker, Helm, Jenkins) on day 3 with HPA autoscaling and a two-stage amd64 image (pymongo's bson C extension SIGILLs on arm64), and fixed a 5-day silent CronJob outage (non-terminal ImagePullBackOff + Forbid concurrency) with paired job and scheduler deadlines.
- Built the FastAPI / MongoDB backend from its first commit to 89 endpoints with ~1,450 offline tests (0.17s) and re-keyed a per-user cache to a shared 10.9 MB catalog (78× smaller) to fit a 512 MiB pod.
- Sole SwiftUI engineer and primary backend author for [SportsBreak](#) (1.4K App Store ratings, 1,119 peak daily active users) and shipped a personalized LLM morning digest to 5,026 users.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine and designed the database schemas and REST APIs for its LangChain LLM features.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer, owning auth, session state, and database interactions, then demoed it until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).
- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers, outperforming Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

SKILLS

Languages

Python, Java, Go, C, SQL, Bash, TypeScript, Swift

ML & Serving

PyTorch, CUDA, RAG, hybrid retrieval, bge embeddings, LLM tool calling / agents

Infra & Observability

Kubernetes (EKS), Docker, Helm, CronJobs, Jenkins, AWS, GCP Vertex AI, Langfuse

Backend & Data

FastAPI, MongoDB, PostgreSQL, Redis Streams, asyncio, pytest, CI/CD, Linux

PROJECTS

- **my-av:** Developed a PyTorch CNN-GRU for real-time steering-angle and speed prediction from 20-frame (1s) video. Built a data pipeline over 33 hours of Comma2k19 with temporal windowing, trained on local CUDA and GCP Vertex AI, and implemented streaming inference.