

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Foundations of Software Engineering, Distributed Systems, Software Refactoring

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Software Engineering, Operating Systems, Algorithms

WORK EXPERIENCE

NewsBreak

May – August 2026

Software Engineering Intern

Mountain View, CA

- Sole SwiftUI engineer and primary backend author for [SportsBreak](#) on a team of ~6, taking it to 1.4K App Store ratings and 1,119 peak daily active users; authored 344 pull requests across five repositories (340 merged).
- Built the Python FastAPI / MongoDB backend from its first commit to 89 REST endpoints with ~1,450 offline unit tests (1:1 with production lines) behind a ruff / pyright / pytest CI gate.
- Stood up the service on AWS EKS (Docker, Helm, Jenkins) on day 3 with 7 Kubernetes CronJobs, then set up Prometheus metrics and Grafana dashboards; diagnosed a DNS → TLS → admission-policy chain that blocked iOS clients on public networks.
- Diagnosed a 5-day silent CronJob outage (non-terminal ImagePullBackOff plus Forbid concurrency) and fixed it with paired job and scheduler deadlines.
- Diagnosed a 65% failure rate in another team's LLM workers via Langfuse traces as an unbounded reasoning loop and a Helm env shadowing the code default; sized a 2048-token cap against observed p95 and labeled a timeout raise a band-aid with a pre-registered falsification test.
- Moved nine third-party integrations server-side so the iOS app ships zero third-party API credentials; cut a 25–30 request client fan-out to one backend call with a shared cache 78× smaller.
- Designed and shipped a personalized LLM alert pipeline, validated it over 11 days of production shadow (4,100+ logged decisions), and cut filler 34% → 4%.
- Introduced the team's Jira backlog, wrote 34 contract-first API specs and ~10,000 lines of docs, onboarded the second engineer, and reopened two closed security tickets with evidence the fixes were absent.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; designed its database schemas and REST APIs.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing protocols, and scalable code practices.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Sole engineer on a VC-backed startup's LLM mock-interview web app, built from 0→1 in React, Next.js, Supabase, and PostgreSQL; owned auth, session state, and database interactions.
- Iterated AI scoring and UI with clients and demoed the product until three 200–500-person companies signed contracts.

SKILLS

Cloud & Systems	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, CI/CD, Linux, Bash, Prometheus, Grafana
Practices	Agile/Scrum, Unit Testing, pytest, Git, Jira, REST APIs
Languages	Python, Java, Go, C, SQL, TypeScript, JavaScript, Swift
Frameworks & Data	FastAPI, React, Next.js, MongoDB, PostgreSQL, MySQL, Supabase, Firebase
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans:** Building a site that rates LLM coding plans. Tested Z.AI, MiniMax, and Kimi on KV caching, long context, and needle-in-a-haystack to show which nerfed caching or heavily quantized models.