Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Distributed Systems, Machine Learning

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Large Language Models in Practice, Operating Systems, Computer Architecture

WORK EXPERIENCE

NewsBreak (Venture Team) **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Built agentic search over SSE: deterministic RAG on turn 1, then a bounded tool-calling agent (≤3 hops, 5 tools, 45 s) with client-held chat state (no PII at rest), RRF hybrid retrieval, and server-resolved citations; entity resolution took real logged queries from 0% to 100% P@1 at 0.4 ms.
- Decomposed personalized news alerts into a multi-step LLM pipeline on a Kubernetes CronJob (bge embeddings → categorical LLM scoring → per-topic percentile gate → LLM dedupe → push) at ∼1,200 LLM calls/day; cut filler 34% → 4% after 11 days of production shadow.
- Root-caused a 65% failure rate in another team's Redis Streams LLM workers via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways) and a Helm env shadowing the code default.
- Built the Python FastAPI / MongoDB backend from its first commit to 89 endpoints with ∼1,450 offline tests and deployed it on AWS EKS (Docker, Helm, autoscaling 2–4 pods) on day 3 with Prometheus and Grafana.
- Primary backend author for SportsBreak (1.4K App Store ratings); moved nine third-party integrations, including LLM and embedding APIs, server-side so the iOS client ships zero API keys.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a Next.js, FastAPI, and Supabase app on GCP App Engine.
- Designed schemas and REST APIs for LangChain LLM resume parsing, suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer in Next.js, Supabase, and PostgreSQL; owned auth, session state, and database interactions.
- Iterated its AI scoring with recruiters and clients until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

AI/ML	Python, LLM tool calling / agents, MCP, RAG, LangChain, PyTorch, CUDA, hybrid retrieval
Backend	Java, FastAPI, REST APIs, SSE, SQL, PostgreSQL, MySQL, MongoDB, Redis Streams
Infra	AWS, GCP Vertex AI, Kubernetes (EKS), Docker, Helm, CI/CD, Prometheus, Grafana, Linux
Languages	C, Go, TypeScript, JavaScript, Swift, Bash
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av:** Built a PyTorch CNN-GRU that predicts steering and speed from 1 s video, with a data pipeline over 33 hours of Comma2k19; trained on local CUDA and GCP Vertex AI and ran streaming inference.
- **codingplans:** Building a site that rates LLM coding plans. Tested Z.AI, MiniMax, and Kimi on KV caching, long context, and needle-in-a-haystack to show which nerfed caching or heavily quantized models.
- **Wisconsin Autonomous:** Integrated wheel-speed encoders on the competition rover and streamed real-time data to the onboard NVIDIA Jetson.