

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Foundations of Software Engineering, Large Language Model Systems, Distributed Systems

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Software Engineering, Large Language Models in Practice, Algorithms

WORK EXPERIENCE

NewsBreak

May – August 2026

Full-Stack Engineer Intern

Mountain View, CA

- Sole SwiftUI engineer and primary backend author for [SportsBreak](#) on a team of ~6, taking it to 1.4K App Store ratings and 1,119 peak daily active users; shipped Firebase auth (Google, Apple, email), a 72-event analytics catalog, and server-driven UI over 27 widget types so new surfaces ship without an App Store release.
- Built the Python FastAPI / MongoDB backend from its first commit to 89 REST endpoints (129 Pydantic models, 26 collections) with ~1,450 offline unit tests behind a ruff / pyright / pytest CI gate.
- Moved nine third-party integrations server-side so the iOS app ships zero third-party API credentials, and shipped a Stripe checkout fulfilled only by a signature-verified, idempotent webhook.
- Cut a 25–30 request client fan-out to one backend call; wrote tests around the AI agent-written cache, then re-keyed it to a shared cache 78× smaller when pod memory neared its 512 MiB limit.
- Shipped agentic search with RAG, a bounded tool-calling agent, and server-resolved citations the model cannot author; deterministic entity resolution took logged queries from 0% to 100% P@1.
- Designed and shipped a personalized LLM alert pipeline, evaluated it over 11 days of production shadow (4,100+ logged decisions), and cut filler 34% → 4%.
- Shipped a personalized LLM morning digest to 5,026 users (~81% of the base), and re-sourced it onto news highlight clusters so every bullet links to its source article.
- Deployed and owned the service on AWS EKS (Docker, Helm, Jenkins, GitHub Actions) from day 3, with Prometheus and Grafana.
- Introduced the team's Jira backlog, wrote 34 contract-first API specs and ~10,000 lines of docs, onboarded the second engineer, and reopened two closed security tickets with evidence the fixes were absent.

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; broke down user stories, led PR reviews, and mentored interns on Git and testing.
- Designed the database schemas and REST APIs for its LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Sole engineer on a VC-backed startup's LLM mock-interview web app, built from 0→1 in React, Next.js, Supabase, and PostgreSQL; owned auth, session state, and database interactions.
- Worked with recruiters so AI scoring matched real hiring, iterated with clients, and demoed the product until three 200–500-person companies signed contracts.

SKILLS

Languages	TypeScript, JavaScript, Python, Java, SQL, Go, Swift, HTML, CSS
Full Stack	React, Next.js, Express, FastAPI, REST APIs, PostgreSQL, MongoDB, Supabase, Firebase
Cloud & CI/CD	AWS, GCP, Docker, Jenkins, GitHub Actions, Kubernetes (EKS), Helm, pytest, Unit Testing
AI	Claude Code / AI coding tools, LLM tool calling / agents, RAG, LangChain, MCP, Langfuse
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans:** Building a site that rates LLM coding plans. Tested Z.AI, MiniMax, and Kimi on KV caching, long context, and needle-in-a-haystack to show which nerfed caching or heavily quantized models.