

Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · [linkedin.com/in/lordphone-wen](https://www.linkedin.com/in/lordphone-wen) · github.com/lordphone

EDUCATION

Carnegie Mellon University

Expected May 2027

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Distributed Systems, Machine Learning, Foundations of Software Engineering

University of Wisconsin–Madison

May 2024

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Operating Systems, Computer Architecture, Algorithms, Mobile Computing Laboratory

WORK EXPERIENCE

NewsBreak

May – August 2026

Software Engineering Intern

Mountain View, CA

- Built a Python FastAPI / MongoDB service from the first commit to 89 endpoints and 7 Kubernetes CronJobs on AWS EKS, with ~1,450 fully-offline tests (0.17s) behind a ruff / pyright / pytest gate.
- Profiled a cache that could not hit (13.2 MB keyed per user) as pod memory neared its 512 MiB limit; re-keyed it to a shared 10.9 MB catalog (78× smaller) and cut the home-load copy 9.0 ms → 0.89 ms.
- Diagnosed a 5-day silent CronJob outage (non-terminal ImagePullBackOff + Forbid concurrency) and fixed it with paired job and scheduler deadlines; set up Prometheus metrics and Grafana dashboards on EKS.
- Root-caused a 65% failure rate in another team's LLM workers via Langfuse traces: an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways) plus a Helm env shadowing the code default.
- Designed an ingest → embed → LLM-score → dedupe pipeline (~2 min ingest-to-decision, ~1,200 LLM calls/day); validated it on 4,100+ production-shadow decisions before going live, cutting filler 34% → 4%.
- Sole SwiftUI engineer and primary backend author for [SportsBreak](#) (~6,000 installs, 1,119 peak daily active users); 344 PRs across five repos (340 merged).

Haddee Education

June – September 2025

Engineering Lead

San Diego, CA

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; designed the database schemas and REST APIs and led PR reviews.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch

September – December 2023

Full Stack Developer

Beijing, China

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; demoed it until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University

February – June 2026

Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers for long-term safety of stochastic systems; outperformed Adaptive MPC in 3-DOF vehicle simulations with LuGre tire models, including ice.
- Implemented Bayesian estimators for real-time road-friction adaptation and mapped natural-language instructions to formal safety specs with multi-turn LLMs.

SKILLS

Languages Python, C, Go, Java, SQL, Bash, Swift, TypeScript

ML & Edge PyTorch, CUDA, OpenCV, NumPy, NVIDIA Jetson, GCP Vertex AI, Bayesian estimation

Infra & Tools Linux, Docker, Kubernetes (EKS), Helm, Jenkins, Prometheus, Grafana, Redis Streams, Git

Certifications AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av**: Built a PyTorch CNN-GRU that predicts steering angle and speed from 20-frame (1s) video; built the data pipeline over 33 hours of Comma2k19 (shuffling, temporal windowing, 240×320 frames), trained on local CUDA and GCP Vertex AI, and ran streaming inference.
- **Wisconsin Autonomous**: Integrated wheel-speed encoders with voltage-to-frequency converters on the competition rover and streamed real-time data to the onboard NVIDIA Jetson.