

EDUCATION

Carnegie Mellon University <i>M.S. in Software Engineering GPA: 4.0 / 4.0</i> Selected Coursework: Distributed Systems, Large Language Model Systems, Machine Learning	Expected May 2027
University of Wisconsin–Madison <i>B.S. in Computer Science & B.S. in Computer Engineering GPA: 3.72 / 4.0</i> Selected Coursework: Operating Systems, Large Language Models in Practice, Algorithms	May 2024

WORK EXPERIENCE

NewsBreak <i>Software Engineering Intern</i>	May – August 2026 <i>Mountain View, CA</i>
- Shipped an LLM search agent over SSE: turn 1 is deterministic RAG; follow-ups run a bounded tool loop (≤ 3 hops, 45 s wall clock) over 5 whitelisted tools, with server-resolved citations so the model cannot author a URL; a failed hybrid-retrieval leg degrades to the survivors, and only all legs failing returns a 500. - Diagnosed a 65% failure rate in another team’s LLM workers via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways); sized a token cap that truncates 0% of successes. - Built an LLM alert pipeline as a Kubernetes CronJob (~1,200 LLM calls/day), calibrated it 11 days in production shadow on 4,100+ logged decisions, then shipped it live, cutting filler 34% $\rightarrow$ 4%. - Stood up the service on production AWS EKS (Docker, Helm, Jenkins) on day 3; diagnosed a 5-day silent CronJob outage (ImagePullBackOff + Forbid concurrency) and fixed it with paired deadlines. - Redesigned an AI-agent-written cache that could not hit (13.2 MB keyed per user) after pod memory neared its 512 MiB limit; re-keyed it to a shared 10.9 MB catalog, $78\times$ smaller. - Built the FastAPI / MongoDB backend from its first commit to 89 endpoints with ~1,450 fully-offline tests in CI; moved nine third-party integrations server-side so the iOS binary ships zero API credentials. - Sole SwiftUI engineer and primary backend author for SportsBreak (~6,000 installs, 1,119 peak daily active users); 344 PRs across five repos (340 merged).	
Haddee Education <i>Engineering Lead</i>	June – September 2025 <i>San Diego, CA</i>
Led an 8-intern team to ship a React, Next.js, FastAPI, and Supabase web app from 0$\rightarrow$1 on GCP App Engine; designed schemas and REST APIs for LangChain resume parsing, suggestions, and AI job matching.	
ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch <i>Full Stack Developer</i>	September – December 2023 <i>Beijing, China</i>
Built a VC-backed startup’s LLM mock-interview web app from 0$\rightarrow$1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; demoed it until three 200–500-person companies signed contracts.	

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University <i>Graduate Researcher Advisor: Prof. Yorie Nakahira</i>	February – June 2026
Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).	

SKILLS

Languages	Python, Go, TypeScript, Java, C, SQL, Bash, Swift
AI & Agents	LLM tool calling / agents, MCP, RAG, Langfuse, Claude Code / AI coding tools
Cloud & Infra	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, Prometheus, Grafana
Backend	FastAPI, asyncio, REST APIs, SSE, MongoDB, PostgreSQL, Redis Streams, pytest
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- codingplans:** Benchmarking paid LLM coding plans (Z.AI, MiniMax, Kimi) with KV-cache, long-context, and needle-in-a-haystack tests to show which providers nerfed caching or heavily quantized their models.
- ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.