Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Large Language Model Systems, Machine Learning

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Algorithms, Software Engineering, Large Language Models in Practice

WORK EXPERIENCE

NewsBreak **May – August 2026**
Full-Stack Engineer Intern *Mountain View, CA*

- Sole SwiftUI engineer and primary backend author for SportsBreak (1.4K App Store ratings, peak DAU 1,119); authored 344 pull requests across five repositories (340 merged).
- Built the FastAPI / MongoDB backend from its first commit to 89 endpoints and ∼1,450 offline tests; stood up production AWS EKS (Docker, Helm, Jenkins) on day 3 with Prometheus / Grafana.
- Shipped an LLM search agent: RAG on turn 1, a bounded tool-calling loop (≤3 hops, 5 tools) for follow-ups, and server-resolved citations; entity resolution took real logged queries 0% → 100% P@1.
- Shipped a personalized LLM news-alert pipeline (∼1,200 LLM calls/day), calibrated in production shadow on 4,100+ logged decisions with prompt ablations; cut filler 34% → 4%.
- Tested an AI-agent-written cache, then redesigned it as pod memory neared 512 MiB: a zero-hit per-user key became one shared cache, 78× smaller; home-load copy 9.0 ms → 0.89 ms.
- Debugged production: a 5-day silent Kubernetes CronJob outage (ImagePullBackOff + Forbid concurrency), and a 65% failure rate in another team's LLM workers, traced via Langfuse to an unbounded reasoning loop.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team to ship a React, Next.js, FastAPI, and Supabase MVP on GCP App Engine.
- Designed schemas and REST APIs for LangChain resume parsing, suggestions, and AI job matching.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing, and scalable code.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch **September – December 2023**
Full Stack Developer *Beijing, China*

- Built the LLM mock-interview app from 0→1 as sole engineer in React, Next.js, and Supabase; owned auth, session state, and database interactions.
- Demoed the product and iterated with clients until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).
- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

SKILLS

Languages & AI	Python, JavaScript, Java, LLM tool calling / agents, RAG, Claude Code / AI coding tools
Full-Stack	TypeScript, React, Next.js, FastAPI, REST APIs, SQL, PostgreSQL, MongoDB, Git
Cloud	AWS, GCP, Kubernetes (EKS), Docker, Helm, CI/CD, GitHub Actions, Grafana
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av:** Built a CNN-GRU model that predicts steering angle and speed from 20-frame (1s) video; processed 33 hours of Comma2k19 footage and ran streaming inference on local CUDA and GCP Vertex AI.
- **ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.