Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Distributed Systems, Machine Learning

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Computer Architecture, Algorithms

WORK EXPERIENCE

NewsBreak **May – August 2026**
Software Engineering Intern *Mountain View, CA*

- Built the Python FastAPI / MongoDB service from its first commit to 89 endpoints and ∼1,450 fully-offline tests (0.17 s) behind a ruff / pyright / pytest GitHub Actions gate; authored 344 PRs (340 merged).
- Stood up the service on production AWS EKS on day 3 (Docker, Helm, Jenkins): 7 CronJobs, HPA autoscaling (2–4 pods at 70% CPU), and Prometheus metrics with Grafana dashboards on the EKS nodes.
- Replaced a 25–30 request client fan-out with one backend call under a 3 s per-source budget, and re-keyed a per-user cache that could never hit to a shared 10.9 MB catalog (78× smaller) sized for a 512 MiB pod; cut home-load copy time 9.0 ms → 0.89 ms.
- Traced a 65% failure rate in another team's LLM workers via Langfuse to an unbounded reasoning loop and a Helm env shadowing the code default; sized a 2,048-token cap against p95 1,197 tokens.
- Diagnosed a 5-day silent CronJob outage (non-terminal ImagePullBackOff + Forbid concurrency) and fixed it with paired job and scheduler deadlines; traced a DNS → TLS → admission-policy chain blocking iOS clients.
- Shipped an SSE-streaming search agent (0.85 s time to first token) over hybrid retrieval (keyword + vector, fused by RRF) that degrades to the surviving legs when one fails, with a bounded tool loop (≤3 hops, 45 s).
- Built a personalized LLM news-alert pipeline as a Kubernetes CronJob (∼1,200 LLM calls/day, ∼2 min ingest to decision), calibrated it on 4,100+ logged decisions in production shadow, and cut filler 34% → 4%.
- Primary backend author for SportsBreak (1.4K App Store ratings); wrote 34 contract-first API specs and ∼10,000 lines of engineering docs, and moved nine third-party integrations server-side.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a Next.js, FastAPI, and Supabase app on GCP App Engine.
- Designed schemas and REST APIs for LangChain LLM resume parsing, suggestions, and AI job matching.
- Led PR reviews and mentored interns on Git workflows, testing protocols, and scalable code practices.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch **September – December 2023**
Full Stack Developer *Beijing, China*

- Built the LLM mock-interview app from 0→1 as sole engineer in React, Next.js, and Supabase; owned auth, session state, and database interactions.
- Iterated AI scoring and UI with users and recruiters; three 200–500-person companies signed contracts.

SKILLS

Languages	Python, Go, C, SQL, Bash, TypeScript, Java, Git
Infra	Linux, Docker, Kubernetes (EKS), Helm, Prometheus, Grafana, GitHub Actions, AWS, GCP
ML & Backend	PyTorch, CUDA, GCP Vertex AI, Langfuse, RAG, FastAPI, asyncio, REST APIs, SSE
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 20-frame (1 s) video; processed 33 hours of Comma2k19 footage and ran streaming inference on local CUDA and GCP Vertex AI.
- **codingplans:** Building a website that rates LLM coding plans; bought plans from Z.AI, MiniMax, and Kimi and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and quantized models.
- **OS coursework:** Implemented a pthread-based quicksort in C and contributed xv6 system-call entry points and a clone wrapper with pointer-argument validation.