Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Data Science for Software Engineering, Machine Learning, Distributed Systems

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Large Language Models in Practice, Software Engineering, Algorithms

WORK EXPERIENCE

NewsBreak **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Built an SSE search agent that splits deterministic from model-driven work: code-planned RAG on turn 1, a bounded tool agent (≤3 hops, 5 whitelisted tools, 45 s) for follow-ups; bare lookups skip the LLM.
- Server-resolved citations so the model cannot author a URL; measured embeddings vs. span matching, then shipped deterministic entity resolution: real logged queries 0% → 100% P@1 at 0.4 ms.
- Built a personalized LLM news-alert pipeline (Kubernetes CronJob, ∼1,200 LLM calls/day) and evaluated it on 4,100+ logged decisions over 11 days of production shadow, hand-graded alerts, a 127-alert holdout, and prompt ablations; cut filler 34% → 4%, with hard caps holding over 11,745 user-days.
- Rejected a rule that caught 91% of filler because it also blocked the day's two highest-reach real alerts; every item logs why it fired or not (`sent`, `below_bar`, `not_event`, ...) for audit.
- Traced a 65% failure rate in another team's LLM workers via Langfuse to an unbounded reasoning loop; sized a 2,048-token cap against p95 1,197 tokens and modeled 49–66% failures vs. 65% observed.
- Built the Python FastAPI / MongoDB backend from its first commit to 89 endpoints and ∼1,450 offline tests behind a ruff / pyright / pytest gate; deployed on AWS EKS (Docker, Helm) with Prometheus / Grafana.
- Primary backend author for SportsBreak (1.4K App Store ratings); moved nine third-party integrations, including LLM APIs, server-side so the client ships zero API keys, and shipped a 72-event analytics catalog.
- Reopened two closed security tickets with line-level evidence that the fixes were absent, and shipped a Stripe checkout whose only trusted fulfillment is a signature-verified, idempotent webhook.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a Next.js, FastAPI, and Supabase app on GCP App Engine.
- Designed schemas and REST APIs for LangChain LLM resume parsing, suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) 2023 Fall Batch **September – December 2023**
Full Stack Developer *Beijing, China*

- Built the LLM mock-interview app from 0→1 as sole engineer in React, Next.js, and Supabase.
- Tuned AI scoring with recruiters to match real hiring; three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

AI & Data	Python, SQL, RAG, LLM tool calling / agents, hybrid retrieval, Langfuse, LangChain, pytest
Backend	FastAPI, Pydantic, REST APIs, SSE, PostgreSQL, MongoDB, TypeScript, React, Next.js
Infra	AWS, GCP, Docker, Kubernetes (EKS), Helm, Terraform, CI/CD, GitHub Actions, Prometheus
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans:** Building a website that rates LLM coding plans; bought plans from Z.AI, MiniMax, and Kimi and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and quantized models.