Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Machine Learning, Large Language Model Systems, Distributed Systems

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Algorithms, Data Structure & Programming III, Large Language Models in Practice

WORK EXPERIENCE

NewsBreak (Venture Team) **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Sole SwiftUI engineer and primary backend author for SportsBreak (1.4K App Store ratings, peak DAU 1,119); authored 344 pull requests across five repositories (340 merged).
- Shipped an LLM search agent in Python: deterministic RAG on turn 1, a bounded tool-calling loop ($\leq$3 hops, 5 tools) for follow-ups, and server-resolved citations; entity resolution took real logged queries 0% $\rightarrow$ 100% P@1.
- Designed and shipped a personalized LLM news-alert pipeline ($\sim$1,200 LLM calls/day); calibrated it in production shadow on 4,100+ logged decisions with prompt ablations, cutting filler 34% $\rightarrow$ 4%.
- Built the FastAPI / MongoDB backend from its first commit to 89 endpoints and $\sim$1,450 offline tests; stood up production AWS EKS (Docker, Helm, Jenkins) on day 3 with Prometheus / Grafana.
- Tested an AI-agent-written cache, then redesigned it as pod memory neared 512 MiB: a zero-hit per-user key became one shared cache, 78$\times$ smaller; home-load copy 9.0 ms $\rightarrow$ 0.89 ms.
- Debugged a 65% failure rate in another team's LLM workers, traced via Langfuse to an unbounded reasoning loop, and a 5-day silent Kubernetes CronJob outage fixed with paired job and scheduler deadlines.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase MVP on GCP App Engine; designed schemas and REST APIs for LangChain resume parsing and AI job matching.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing, and scalable code.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview app from 0$\rightarrow$1 as sole engineer; worked with recruiters so AI scoring categories matched real hiring, and iterated scoring and UI from user-testing sessions.
- Demoed the product and iterated fast with clients until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages & AI	Python, SQL, LLM tool calling / agents, RAG, Claude Code / AI coding tools, PyTorch
Full-Stack	TypeScript, JavaScript, Java, React, Next.js, FastAPI, REST APIs, PostgreSQL, MongoDB
Cloud	AWS, GCP, GCP Vertex AI, Kubernetes (EKS), Docker, Helm, CI/CD, Grafana
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and heavy quantization.
- **my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 1 s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI, with streaming inference.
- **ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.