Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Large Language Model Systems, Foundations of Software Engineering

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Software Engineering, Algorithms

WORK EXPERIENCE

NewsBreak **May – August 2026**
Backend Engineer Intern *Mountain View, CA*

- Built the SportsBreak FastAPI / MongoDB backend from its first commit to 89 REST endpoints, with ∼1,450 offline tests gating every merge (ruff, pyright, pytest); wrote 34 contract-first API specs.
- Shipped Firebase auth (Google, Apple, email) with guest-first sessions upgraded in place, and moved nine third-party integrations server-side so the iOS binary ships zero third-party API credentials.
- Shipped Stripe web checkout for DramaBreak with a signature-verified webhook as the only trusted fulfillment path and idempotent fulfillment that reclaims stale pending payments.
- Built an LLM search API that streams answers over SSE: fixed RAG retrieval on the first question, a capped tool-calling agent on follow-ups, and server-resolved citations the model cannot author.
- Stood up the service on EKS (Helm, Jenkins) on day 3 and set up Prometheus / Grafana; diagnosed a 5-day silent CronJob outage (ImagePullBackOff + Forbid concurrency) and fixed it with paired deadlines.
- Root-caused a 65% failure rate in another team's LLM workers via Langfuse traces to an unbounded reasoning loop and a Helm env shadowing the code default; sized a token cap to the failing path.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; designed database schemas and REST APIs for LangChain resume parsing and AI job matching.
- Led PR reviews, broke down user stories, and mentored interns on Git workflows, testing, and scalable code.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built the LLM mock-interview web app from 0→1 as sole engineer in React, Next.js, and Supabase (PostgreSQL); owned auth, session state, and database interactions.
- Demoed the product to clients and iterated AI scoring with recruiters so its categories matched real-life hiring; three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation.

SKILLS

Languages	Python, Go, JavaScript, TypeScript, SQL, Bash, Swift, Java
APIs & Practices	REST APIs, FastAPI, Server-Sent Events, Stripe, Git, GitHub Actions, CI/CD, pytest
Cloud & Infra	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, Prometheus, Grafana, Linux
Data & AI	MongoDB, PostgreSQL, Supabase, React, Next.js, LangChain, RAG, LLM tool calling
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **codingplans:** Building a site that rates LLM coding plans; bought plans from Z.AI, MiniMax, and Kimi and tested KV caching and long context to show which providers nerfed caching or quantized models.
- **my-av:** Built a CNN-GRU that predicts steering and speed from 1s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI with streaming inference.