Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Distributed Systems, Machine Learning

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Large Language Models in Practice, Computer Architecture

WORK EXPERIENCE

NewsBreak **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Designed and shipped a personalized LLM alert pipeline as a Kubernetes CronJob: bge embeddings shortlist near-duplicates, an LLM scores each article, a per-topic 7-day percentile gate selects, and an LLM same-story judge dedupes before APNs (∼2 min ingest-to-decision).
- Logged why each item fired or not and calibrated in production shadow (4,100+ decisions, 127-alert holdout, prompt ablations); live: 1.8 median alerts/user/day at 4% filler, 2/day cap held over 11,745 user-days.
- Designed an SSE search agent: deterministic RAG on turn 1, then a bounded tool-calling loop (≤3 hops, 5 tools, 45 s) over hybrid retrieval (keyword + bge vectors, fused by RRF) with server-resolved citations; entity resolution reached 100% P@1 on real logged queries at 0.4 ms.
- Root-caused a 65% failure rate in another team's LLM workers via Langfuse traces to an unbounded reasoning loop and a Helm env shadowing the code default; sized a 2,048-token cap against p95 1,197 tokens.
- Built the SportsBreak (1.4K App Store ratings) FastAPI / MongoDB backend from its first commit to 89 endpoints and ∼1,450 fully-offline tests; stood it up on AWS EKS on day 3 and set up Prometheus / Grafana.
- Root-caused a 5-day silent Kubernetes CronJob outage: an empty image tag left the Job in non-terminal ImagePullBackOff, so Forbid concurrency skipped every later run; fixed it with paired deadlines.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a Next.js, FastAPI, and Supabase app on GCP App Engine; designed schemas and REST APIs for LangChain resume parsing and AI job matching.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing, and scalable code.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built the LLM mock-interview app from 0→1 as sole engineer in React, Next.js, and Supabase; owned auth, session state, and database interactions.
- Iterated AI scoring with recruiters to match real hiring; three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

AI & Agents	Python, LLM tool calling / agents, RAG, RRF, bge embeddings, MCP, Claude Code
Infra	Kubernetes (EKS), Docker, Helm, Prometheus, Grafana, Langfuse, AWS, Linux, CI/CD
Data & Backend	C, Go, SQL, MongoDB, PostgreSQL, Redis Streams, FastAPI, asyncio, pytest, LangChain
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **Operating Systems (C):** Implemented a pthread-based quicksort for UW–Madison coursework and contributed xv6 system-call entry points and a clone wrapper with pointer-argument validation.
- **ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.