Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Machine Learning, Distributed Systems

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Computer Architecture, Operating Systems, Algorithms

WORK EXPERIENCE

NewsBreak **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Diagnosed a 65% failure rate in another team's LLM workers via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. runaways at 16,384) and sized a 2,048-token cap.
- Traced pod memory climbing toward its 512 MiB limit to a cache that could never hit (13.2 MB per user); re-keyed it to one shared 10.9 MB copy and cut home-load copy time from 9.0 ms to 0.89 ms.
- Built an LLM news-alert pipeline (embeddings, DeepSeek scoring, LLM dedupe) at ∼1,200 calls/day and ∼2 min ingest-to-decision; shadow-tested it 11 days on production traffic, cutting filler 34% → 4%.
- Built an LLM search API that streams answers over SSE (0.85 s to first token): hybrid vector + text retrieval fused by RRF, fixed RAG on the first question, and a capped tool-calling agent on follow-ups.
- Replaced embeddings with deterministic entity resolution after measuring both (0.4 ms per resolve); got 20/20 real logged search queries right that used to return nothing.
- Built the SportsBreak FastAPI / MongoDB backend from its first commit to 89 endpoints on AWS EKS with ∼1,450 offline tests gating CI, and set up Prometheus metrics and Grafana dashboards.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; designed the schemas and REST APIs for LangChain resume parsing and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built the LLM mock-interview web app from 0→1 as sole engineer and iterated its AI scoring with recruiters and clients until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages	Python, C, Bash, SQL, Go, Java, TypeScript, Swift
ML & Inference	PyTorch, CUDA, NumPy, GCP Vertex AI, DeepSeek / DeepInfra, RAG, Langfuse
Systems & Infra	Linux, AWS, GCP, Kubernetes (EKS), Docker, Prometheus, Grafana, asyncio
Backend	FastAPI, MongoDB, PostgreSQL, REST APIs, Server-Sent Events, CI/CD, pytest, Git
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **codingplans:** Building a site that rates LLM coding plans; bought plans from Z.AI, MiniMax, and Kimi and tested KV caching, long context, and needle-in-a-haystack to show which providers nerfed caching or heavily quantized models.
- **my-av:** Built a PyTorch CNN-GRU that predicts steering and speed from 1s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI, with streaming inference.
- **OS coursework:** Implemented a pthread-based quicksort in C and contributed xv6 system-call entry points and a clone wrapper with pointer-argument validation.