Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Machine Learning, Large Language Model Systems, Distributed Systems

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Large Language Models in Practice, Computer Architecture

WORK EXPERIENCE

NewsBreak **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Designed and deployed a personalized LLM alert pipeline as a Kubernetes CronJob: bge embeddings, per-article LLM scoring, a per-topic 7-day percentile gate, and an LLM same-story judge (∼1,200 LLM calls/day).
- Evaluated it in production shadow (4,100+ logged decisions, a 127-alert holdout, prompt ablations) before going live; cut filler 34% → 4% at 1.8 median sends per user per day.
- Designed an SSE search agent: deterministic RAG, then a bounded tool-calling loop (≤3 hops, 5 tools) over hybrid retrieval fused by RRF; entity resolution hit 100% P@1 on real logged queries at 0.4 ms.
- Root-caused a 65% failure rate in another team's LLM workers via Langfuse traces to an unbounded reasoning loop and a Helm env shadowing the code default; sized a 2,048-token cap against p95 1,197 tokens.
- Stood up production AWS EKS on day 3 (Docker, Helm, Jenkins), set up Prometheus and Grafana monitoring, and fixed a 5-day silent CronJob outage with paired job and scheduler deadlines.
- Built the SportsBreak (1.4K App Store ratings) FastAPI / MongoDB backend from its first commit to 89 endpoints; re-keyed a cache that could never hit into a shared one, 78× smaller, to fit a 512 MiB pod.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints and PR reviews to ship a Next.js, FastAPI, and Supabase app on GCP App Engine; designed REST APIs for LangChain LLM resume parsing and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview app from 0→1 as sole engineer in React, Next.js, Supabase, and PostgreSQL; iterated AI scoring with recruiters until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).
- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

SKILLS

ML & GenAI	Python, PyTorch, CUDA, NumPy, LLM agents, RAG, bge embeddings, Langfuse
Cloud & Infra	AWS, Kubernetes (EKS), Docker, Helm, Prometheus, Grafana, GCP Vertex AI, Linux
Data & Systems	SQL, C, Go, PostgreSQL, MongoDB, Redis Streams, FastAPI, asyncio, pytest, Git
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av:** PyTorch CNN-GRU predicting steering and speed from 1 s video; trained on 33 hours of Comma2k19 on CUDA and GCP Vertex AI.
- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and heavy quantization.
- **Operating systems (C):** Implemented a pthread-based quicksort in C and contributed xv6 system-call entry points and a clone wrapper.