Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Machine Learning, Distributed Systems

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Large Language Models in Practice, Computer Architecture, Operating Systems

WORK EXPERIENCE

NewsBreak (Venture Team) **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Built an LLM search agent: fixed RAG on the first question, a capped tool-calling loop (5 tools, ≤3 hops) on follow-ups, and citations the model cannot author; got 20/20 logged queries right that used to return nothing.
- Built a personalized breaking-news pipeline (embeddings, LLM scoring, LLM same-story judge); hand-graded alerts, ran prompt ablations, and shadowed 11 days of production traffic before launch; filler fell 34% → 4%.
- Killed my own first alert detector two days after shipping it when 19 hours of production data showed it was 15 hours too slow, deleting the code and its calibration harness (net −882 lines).
- Diagnosed a 65% failure rate in another team's LLM workers: Langfuse traces showed an unbounded reasoning loop hitting a 16,384-token ceiling, so I sized a 2,048-token cap against p95 1,197 tokens.
- Shipped SportsBreak 0→1 as sole SwiftUI engineer and built its FastAPI / MongoDB backend from the first commit to 89 endpoints on AWS EKS, with ∼1,450 offline tests; its LLM morning digest reached 81% of users.
- Started the Jira backlog unprompted, wrote 34 contract-first API specs, and onboarded the second engineer; worked with legal to get through three App Store trademark rejections.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; designed schemas and REST APIs for LangChain resume parsing and AI job matching.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows, testing, and scalable code.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer in React, Next.js, and Supabase; owned auth, session state, and database interactions.
- Iterated AI scoring and UI with clients and recruiters until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

AI & Agents	LLM tool calling / agents, MCP, RAG, hybrid retrieval, Langfuse, LangChain, Claude Code
Full Stack	Python, TypeScript, Swift, SQL, FastAPI, React, Next.js, SwiftUI, MongoDB, PostgreSQL
Infra	AWS, Kubernetes (EKS), Docker, Helm, GitHub Actions, Prometheus, Grafana, GCP, Linux
Practices	REST APIs, Server-Sent Events (SSE), CI/CD, pytest, Git, Agile/Scrum, Jira
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **codingplans:** Building a website that rates LLM coding plans; bought plans from Z.AI, MiniMax, and Kimi and tested KV caching, long context, and needle-in-a-haystack to show which providers nerfed caching or heavily quantized models.
- **ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.