Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Data Science for Software Engineering, Machine Learning

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Algorithms, Operating Systems, Software Engineering, Data Structure & Programming III

WORK EXPERIENCE

NewsBreak **May – August 2026**
Software Engineering Intern *Mountain View, CA*

- Built a near-real-time news pipeline as a Kubernetes CronJob (ingest, embed, LLM-score, dedupe, push in ∼2 min); ran it in shadow on production traffic for 11 days, cutting filler alerts from 34% to 4%.
- Built the Python FastAPI and MongoDB backend from its first commit to 89 REST endpoints and 7 Kubernetes CronJobs on AWS EKS, with ∼1,450 offline tests gating every merge in CI (ruff, pyright, pytest).
- Stood up production EKS (Helm, Jenkins, Docker) on day 3, set up Prometheus and Grafana, and fixed a 5-day silent CronJob outage with paired job and scheduler deadlines.
- Cut the World feed from 25–30 client requests to one backend call, and fixed a per-user cache that never hit (13.2 MB each) with one shared 10.9 MB copy as the pod neared its memory limit.
- Diagnosed a 65% failure rate in another team's LLM pipeline, via Langfuse traces, as an unbounded reasoning loop hitting a 16,384-token ceiling, and sized a token cap against observed successes.
- Shipped SportsBreak to the App Store as the sole SwiftUI engineer (∼6,000 installs); migrated its ads to AppLovin MAX and grew `app-ads.txt` from 62 to 1,332 records, recovering three days of lost bids.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase MVP on GCP App Engine; designed database schemas and REST APIs for LangChain resume parsing and AI job matching.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows and testing protocols.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned authentication, session state, and database interactions.
- Iterated AI scoring and UI from user testing, worked with recruiters so scoring matched real hiring, and demoed the product until three 200–500-person companies signed contracts.

SKILLS

Languages	Python, Java, SQL, TypeScript, JavaScript, Go, C, Swift
Cloud & Data	AWS, Terraform, Docker, Kubernetes (EKS), PostgreSQL, MySQL, MongoDB, Redis Streams
DevOps	Helm, Jenkins, GitHub Actions, CI/CD, CronJobs, Prometheus, Grafana, Linux, GCP
AI / LLM	Claude Code / AI coding tools, RAG, LLM tool calling / agents, LangChain, Langfuse, PyTorch
Practices	FastAPI, REST APIs, pytest, Unit Testing, TDD, Git, Agile/Scrum, Jira
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **my-av:** Developed a PyTorch CNN-GRU for real-time steering-angle and speed prediction from 20-frame (1s) video; built a data pipeline over 33 hours of Comma2k19 (shuffling, temporal windowing, 240×320 frames).
- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and heavy quantization.
- **Wisconsin Autonomous:** Integrated wheel-speed encoders with voltage-to-frequency converters on the competition rover and streamed real-time data to the onboard NVIDIA Jetson.
- **ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.