Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Machine Learning, Large Language Model Systems, Data Science for Software Engineering

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Large Language Models in Practice, Algorithms, Data Structure & Programming III

WORK EXPERIENCE

NewsBreak **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Built a RAG search agent over hybrid retrieval (keyword, Mongo text search, and bge embeddings fused by RRF), with a capped tool-calling loop on follow-ups and server-resolved citations the model cannot author.
- Benchmarked embeddings against span-based entity resolution on real logged queries, rejected embeddings with measurements, and took queries that returned nothing to 20/20 correct in an eval harness on the real catalog.
- Built a personalized news-alert pipeline (bge embeddings, LLM scoring into an event category and 0–100 importance, LLM duplicate check); hand-graded alerts, ran prompt ablations in shadow mode on 11 days of production traffic, cutting filler alerts from 34% to 4%.
- Diagnosed a 65% failure rate in another team's LLM pipeline, via Langfuse traces, as an unbounded reasoning loop hitting a 16,384-token ceiling, and sized a 2,048-token cap against p95 output of 1,197 tokens.
- Built the Python FastAPI and MongoDB backend from its first commit to 89 REST endpoints on AWS EKS (Docker, Helm), with ∼1,450 offline tests gating CI and Prometheus / Grafana monitoring.
- Shipped SportsBreak to the App Store as sole SwiftUI engineer and primary backend author (∼6,000 installs, 1,119 peak DAU), plus a personalized LLM morning digest that reached 5,026 users (∼81% of the base).

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; designed database schemas and REST APIs for LangChain LLM features, including resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer (React, Next.js, Supabase, PostgreSQL); worked with recruiters so AI scoring categories matched real hiring, and iterated with clients until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).
- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

PROJECTS

- **my-av:** Developed a PyTorch CNN-GRU for real-time steering-angle and speed prediction from 20-frame (1s) video; built a data pipeline over 33 hours of Comma2k19 with shuffling, temporal windowing, and 240×320 frame processing. Trained on local CUDA and GCP Vertex AI and implemented streaming inference.

SKILLS

Languages & ML	Python, SQL, PyTorch, CUDA, NumPy, OpenCV, Bayesian estimation
NLP & LLMs	RAG, hybrid retrieval, bge embeddings, LLM tool calling / agents, LangChain, Langfuse
Cloud & Data	AWS, GCP Vertex AI, Docker, Kubernetes (EKS), FastAPI, MongoDB, PostgreSQL
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner