Lordphone Wen

Toronto, ON

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Machine Learning, Large Language Model Systems, Data Science for Software Engineering, Distributed Systems

University of Wisconsin–Madison **May 2024**

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Large Language Models in Practice, Algorithms, Operating Systems

WORK EXPERIENCE

NewsBreak **May – August 2026**

AI Engineer Intern *Mountain View, CA*

- Designed and shipped a personalized LLM data pipeline as a Kubernetes CronJob (ingest → embed → LLM-classify and score → per-topic percentile select → LLM same-story dedupe → push in ∼2 min; ∼1,200 LLM calls/day), logging why each item fired or did not.
- Evaluated it over 11 days of production shadow on 4,100+ logged decisions: hand-graded alerts showed filler and real events both scored 89–90, so added a categorical *happened / speculative / no_event* gate, checked it on a 127-alert holdout, and ran prompt ablations; rejected a rule that caught 91% of filler but blocked the day's two highest-reach real alerts; shipped live with filler 34% → 4%.
- Shipped RAG and agentic search on one SSE endpoint: turn 1 is deterministic RAG over hybrid retrieval (substring + text + vector embeddings, fused by RRF); follow-ups run a bounded tool-calling agent ($\leq$3 hops, 5 tools) with server-resolved citations the model cannot author.
- Built an eval harness on real logged queries against the production catalog; deterministic entity resolution answered 20/20 queries that used to return nothing, at 0.4 ms, after benchmarks ruled out embeddings.
- Diagnosed a 65% failure rate in another team's LLM pipeline via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. 16,384 runaways) and sized a 2,048-token cap against observed successes.
- Sole backend author and primary iOS engineer for SportsBreak (∼6,000 installs): built the FastAPI / MongoDB backend from the first commit to 89 endpoints and 7 CronJobs on AWS EKS (Docker, Helm), with ∼1,450 offline tests gating CI; set up Prometheus and Grafana dashboards.

Haddee Education **June – September 2025**

Engineering Lead *San Diego, CA*

- Designed database schemas and REST APIs for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.
- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**

Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase, PostgreSQL); worked with recruiters so AI scoring categories matched real hiring, and iterated with clients until three 200–500-person companies signed contracts.

SKILLS

AI & ML	RAG, hybrid retrieval (RRF), LLM tool calling / agents, LangChain, PyTorch, Langfuse
Languages	Python, SQL, Go, Java, TypeScript, C
Cloud & MLOps	AWS, GCP Vertex AI, Kubernetes (EKS), Docker, Helm, CI/CD, Prometheus, Grafana
Data & Backend	MongoDB, PostgreSQL, MySQL, Redis Streams, FastAPI, Pydantic, REST APIs, pytest
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 1 s video; built a data pipeline over 33 hours of Comma2k19 and trained on local CUDA and GCP Vertex AI, with streaming inference.
- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and benchmarked KV caching, long context, and needle-in-a-haystack to expose nerfed caching and quantization.