Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**

M.S. in Software Engineering | GPA: 4.0 / 4.0

Selected Coursework: Distributed Systems, Data Science for Software Engineering, Machine Learning, Large Language Model Systems

University of Wisconsin–Madison **May 2024**

B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0

Selected Coursework: Algorithms, Data Structure & Programming III, Software Engineering

WORK EXPERIENCE

NewsBreak **May – August 2026**

Full-Stack Engineer Intern *Mountain View, CA*

- Shipped SportsBreak to the App Store as primary iOS engineer and sole backend author (∼6,000 installs, 1,119 peak daily active users); built the FastAPI / MongoDB backend from its first commit to 89 REST endpoints and 7 Kubernetes CronJobs, with ∼1,450 offline tests gating every merge in CI (ruff, pyright, pytest).
- Stood up the service on production AWS EKS (Docker, Helm, Jenkins) on day 3, diagnosed a DNS → TLS → admission-policy chain blocking iOS clients on public networks, and set up Prometheus / Grafana.
- Designed a personalized LLM data pipeline as a Kubernetes CronJob (ingest → embed → LLM-score → select → LLM same-story dedupe → push in ∼2 min); ran it in shadow on production traffic for 11 days before launch, cutting filler alerts from 34% to 4%.
- Built an LLM search API: deterministic RAG over hybrid retrieval on the first question, a bounded tool-calling agent on follow-ups, and server-resolved citations; fixed 20/20 logged queries that returned nothing.
- Built server-driven UI over 27 widget types so new surfaces ship without an App Store release, and migrated nine third-party integrations server-side so the app ships zero third-party API credentials.
- Introduced the team's Jira backlog, wrote 34 contract-first API specs and ∼10,000 lines of engineering docs, onboarded the second engineer, and worked with legal on launch guidance.

Haddee Education **June – September 2025**

Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 and ship it on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.
- Designed database schemas and REST APIs for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**

Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase, PostgreSQL); demoed it and iterated with clients until three 200–500-person companies signed contracts.

Control & Learning Group — Carnegie Mellon University **February – June 2026**

Graduate Researcher

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation.

SKILLS

Languages	Python, TypeScript, JavaScript, SQL, Swift, Go, Java, HTML, CSS
Full Stack	FastAPI, React, Next.js, REST APIs, Pydantic, Supabase, Firebase, pytest
Cloud & CI/CD	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, CI/CD, Grafana
Data & AI	MongoDB, PostgreSQL, MySQL, RAG, LLM tool calling / agents, LangChain, PyTorch
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 1 s video; built a data pipeline over 33 hours of Comma2k19 and trained on local CUDA and GCP Vertex AI, with streaming inference.
- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and benchmarked KV caching, long context, and needle-in-a-haystack to expose nerfed caching.