Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Foundations of Software Engineering, Software Refactoring

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Data Structure & Programming III, Computer Architecture

WORK EXPERIENCE

NewsBreak **May – August 2026**
Backend Engineer Intern *Mountain View, CA*

- Shipped SportsBreak to the App Store as sole backend author and primary iOS engineer; built the FastAPI / MongoDB service from its first commit to 89 REST endpoints and 7 Kubernetes CronJobs, with ∼1,450 offline tests gating every merge in CI (ruff, pyright, pytest).
- Wrote 34 contract-first API specs and ∼10,000 lines of docs; built server-driven UI over 27 widget types so new surfaces ship without an App Store release and unknown kinds drop instead of failing.
- Stood up the service on production AWS EKS (Docker, Helm, Jenkins) on day 3; debugged a DNS → TLS → admission-policy chain that blocked iOS clients on public networks, and fixed a 5-day silent CronJob outage (non-terminal ImagePullBackOff + Forbid concurrency) with paired job and scheduler deadlines.
- Migrated nine third-party integrations server-side so the app ships zero third-party API credentials; found that iOS silently drops a Live Activity start push with no `aps.alert` while APNs returns HTTP 200.
- Built a personalized LLM alert pipeline as an automated CronJob (ingest → embed → score → dedupe → push); ran it in shadow on production traffic for 11 days before launch, cutting filler alerts from 34% to 4%.
- Built a streaming (SSE) LLM search API with a bounded tool-calling agent and server-resolved citations; retrieval legs degrade independently, and entity resolution fixed 20/20 logged queries that returned nothing.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 and ship it on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.
- Designed and reviewed the database schemas and REST API structure for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase, PostgreSQL); owned auth and the database and iterated with clients until three 200–500-person companies signed.

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher

- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty, including ice.

SKILLS

Languages	Python, Go, C, Java, TypeScript, Swift, SQL, Bash
APIs & Backend	REST APIs, Server-Sent Events, FastAPI, Pydantic, asyncio, Redis Streams, pytest, Git
Cloud & DevOps	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, GitHub Actions, CI/CD, Prometheus, Grafana
Data & AI	MongoDB, PostgreSQL, RAG, LLM tool calling / agents, LangChain, PyTorch
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **Wisconsin Autonomous:** Integrated wheel-speed encoders with voltage-to-frequency converters on the competition rover and streamed real-time data to the onboard NVIDIA Jetson.
- **my-av:** Built a PyTorch CNN-GRU that predicts steering angle and speed from 1 s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI, with streaming inference.