Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Machine Learning, Large Language Model Systems

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Algorithms, Computer Architecture

WORK EXPERIENCE

NewsBreak **May – August 2026**
Software Engineering Intern *Mountain View, CA*

- Built an LLM search API over hybrid retrieval (substring + Mongo text + bge vector, fused by RRF) that degrades to the surviving legs when one fails; deterministic RAG on the first question, a capped tool-calling agent on follow-ups, and server-resolved citations streamed over SSE.
- Made deterministic entity resolution the query rewriter (0.4 ms per resolve), getting 20/20 real logged queries right that used to return nothing; rejected embeddings with measurements.
- Designed an ingest → embed → LLM-score → dedupe → push pipeline as a Kubernetes CronJob (∼2 min ingest-to-decision); ran it in shadow on production traffic for 11 days, cutting filler alerts from 34% to 4%.
- Shipped SportsBreak as sole backend author and primary iOS engineer; built the Python FastAPI / MongoDB service from its first commit to 89 REST endpoints, with ∼1,450 offline tests gating every merge in CI.
- Stood up the service on production AWS EKS (Docker, Helm, Jenkins) on day 3 with HPA autoscaling and Prometheus / Grafana monitoring; root-caused a 5-day silent CronJob outage and a 65% failure rate in another team's LLM pipeline (via Langfuse traces).
- Fixed a cache that never hit because it was keyed per user (13.2 MB each) by sharing one 10.9 MB copy as the pod neared its 512 MiB memory limit, and cut a feed from 25–30 client requests to one backend call.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 and ship it on GCP App Engine; led PR reviews and mentored interns on Git workflows, testing protocols, and scalable code.
- Designed and reviewed the database schemas and REST API structure for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase, PostgreSQL); worked with recruiters and clients on requirements until three 200–500-person companies signed.

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages	Python, Java, Go, JavaScript, TypeScript, C, SQL, Bash
Search & ML	RAG, hybrid retrieval, RRF, bge embeddings, PyTorch, CUDA, LLM tool calling / agents
Cloud & CI/CD	GCP, AWS, Kubernetes (EKS), Docker, Linux, CI/CD, GitHub Actions, Jenkins
Backend & Web	REST APIs, FastAPI, MongoDB, PostgreSQL, Redis Streams, pytest, React, Next.js
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **my-av:** Built a PyTorch CNN-GRU that predicts steering and speed from 1 s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI with streaming inference.
- **Operating Systems (UW–Madison):** Implemented a pthread-based quicksort in C and contributed xv6 system-call entry points and a clone wrapper with pointer-argument validation.