Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Machine Learning, Foundations of Software Engineering

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Large Language Models in Practice, Software Engineering, Algorithms

WORK EXPERIENCE

NewsBreak (Venture Team) **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Shipped an LLM search assistant on one SSE endpoint: RAG over hybrid retrieval on the first question, then a bounded tool-calling agent (≤3 hops, 5 tools) with server-resolved citations the model cannot author; bare entity lookups skip the LLM.
- Built an eval harness on real logged queries against the production catalog; new entity resolution got 20/20 queries right that used to return nothing, after measurements ruled out embeddings.
- Hand-graded a personalized LLM alert pipeline, found filler and real events both scored 89–90, and added a categorical gate; ran it in production shadow for 11 days before launch, cutting filler 34% → 4%.
- Primary iOS engineer and sole backend author for SportsBreak (∼6,000 installs, 1,119 peak DAU): built the FastAPI / MongoDB backend from its first commit to 89 endpoints on AWS EKS with ∼1,450 offline tests gating CI, and shipped an LLM morning digest to 5,026 users (∼81% of the base).
- Shipped Stripe web checkout on DramaBreak with a signature-verified, idempotent webhook; worked with legal through three App Store trademark rejections, cutting the hero World Cup feature to launch.
- Diagnosed a 65% failure rate in another team's LLM pipeline via Langfuse traces as an unbounded reasoning loop (p95 1,197 tokens vs. runaways at 16,384) and sized a 2,048-token cap.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Designed schemas and REST APIs for LangChain resume parsing, suggestions, and AI job matching.
- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as sole engineer in React, Next.js, and Supabase (PostgreSQL); owned auth, session state, and database interactions.
- Worked with recruiters so AI scoring matched real hiring, iterated on scoring and UI from user-testing sessions, and demoed to clients until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages	TypeScript, JavaScript, Python, SQL, Go, Java, Swift
AI & LLM	LLM tool calling / agents, RAG, hybrid retrieval (RRF), LangChain, Langfuse, Claude Code
Full-Stack	React, Next.js, Tailwind CSS, FastAPI, Pydantic, REST APIs, Server-Sent Events, Stripe
Data & Cloud	PostgreSQL, MongoDB, Supabase, AWS, Kubernetes (EKS), Docker, pytest, CI/CD
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching and quantization.