Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Foundations of Software Engineering, Large Language Model Systems

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Computer Architecture, Algorithms

WORK EXPERIENCE

NewsBreak **May – August 2026**
Backend Engineer Intern *Mountain View, CA*

- Shipped SportsBreak to the App Store as sole backend author and primary iOS engineer; built the Python FastAPI / MongoDB service from its first commit to 89 REST endpoints and 7 Kubernetes CronJobs, with ∼1,450 offline tests gating every merge in CI (ruff, pyright, pytest).
- Stood up the containerized service on production AWS EKS (Docker, Helm, Jenkins) on day 3 and set up Prometheus metrics and Grafana dashboards; debugged a DNS → TLS → admission-policy chain that blocked iOS clients on public networks and drove a public-domain fix.
- Root-caused a 5-day silent CronJob outage (non-terminal ImagePullBackOff + Forbid concurrency) and fixed it with paired job and scheduler deadlines; diagnosed a 65% failure rate in another team's LLM pipeline, via Langfuse traces, as an unbounded reasoning loop hitting a 16,384-token ceiling.
- Fixed a cache that never hit because it was keyed per user (13.2 MB each) by sharing one 10.9 MB copy as the pod neared its 512 MiB memory limit, and cut the World feed from 25–30 client requests to one backend call.
- Migrated nine third-party integrations server-side so the app ships zero third-party API credentials; wrote 34 contract-first API specs and ∼10,000 lines of docs, and onboarded the team's second engineer.
- Built a personalized LLM alert pipeline as a Kubernetes CronJob (ingest → embed → score → dedupe → push); ran it in shadow on production traffic for 11 days before launch, cutting filler alerts from 34% to 4%.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 and ship it on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.
- Designed and reviewed the database schemas and REST API structure for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase, PostgreSQL); owned auth and the database and iterated with clients until three 200–500-person companies signed.

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher

- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty, including ice.

SKILLS

Languages	Python, Go, C, TypeScript, JavaScript, Java, Swift, SQL, Bash
Cloud & DevOps	AWS, GCP, Kubernetes (EKS), Docker, Helm, GitHub Actions, Prometheus, Grafana
Backend & Testing	REST APIs, FastAPI, Pydantic, asyncio, pytest, Unit Testing, Redis Streams, Git
Frontend & Data	React, Next.js, MongoDB, LLM tool calling / agents, Claude Code / AI coding tools
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **Operating Systems (UW–Madison):** Implemented a pthread-based quicksort in C and contributed xv6 system-call entry points and a clone wrapper with pointer-argument validation.
- **my-av:** Built a CNN-GRU that predicts steering and speed from 1 s video; trained on 33 hours of Comma2k19 on local CUDA and GCP Vertex AI.