Lordphone Wen

New York, NY

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Distributed Systems, Machine Learning

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Large Language Models in Practice, Operating Systems, Algorithms

WORK EXPERIENCE

NewsBreak **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Shipped SportsBreak to the App Store as the team's primary iOS engineer and sole backend author; the app reached ∼6,000 installs and 1,119 peak daily active users.
- Built the Python FastAPI and MongoDB backend from its first commit to 89 REST endpoints and 7 Kubernetes CronJobs on AWS EKS (Docker, Helm, Jenkins), with ∼1,450 offline tests gating CI.
- Built an LLM search agent: hybrid RAG (keyword + bge vectors, RRF) on first questions, a capped tool-calling loop (5 tools, ≤3 hops) on follow-ups, SSE streaming, and server-resolved citations; an eval harness took logged queries that returned nothing to 20/20 correct.
- Built a personalized news-alert pipeline (embeddings, LLM scoring, LLM duplicate check) as a Kubernetes CronJob; ran prompt ablations in shadow mode on 11 days of production traffic, cutting filler from 34% to 4%.
- Diagnosed a 65% failure rate in another team's LLM pipeline, via Langfuse traces, as an unbounded reasoning loop hitting a 16,384-token ceiling, and sized a 2,048-token cap against p95 output of 1,197 tokens.
- Redesigned a cache as pod memory neared its 512 MiB limit: keyed per user, it never hit and held 13.2 MB per user, so re-keyed it on content to one shared 10.9 MB copy.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 on GCP App Engine; designed schemas and REST APIs for LangChain resume parsing and AI job matching.
- Broke down user stories, led PR reviews, and mentored interns on Git workflows and testing.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Built a VC-backed startup's LLM mock-interview web app from 0→1 as the sole engineer in React, Next.js, Supabase, and PostgreSQL; owned authentication, session state, and database interactions.
- Iterated AI scoring and UI from user testing and fast client feedback, working with recruiters so scoring matched real hiring, until three 200–500-person companies signed contracts.

RESEARCH EXPERIENCE

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher | Advisor: Prof. Yorie Nakahira

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages & AI	Python, Go, SQL, LLM tool calling / agents, MCP (Model Context Protocol), RAG
Cloud & Infra	AWS, Kubernetes (EKS), Docker, Helm, CI/CD, Prometheus, Grafana, Git
Backend	FastAPI, Pydantic, asyncio, REST APIs, LangChain, MongoDB, PostgreSQL, Redis Streams
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS & LEADERSHIP

- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and tested KV caching, long context, and needle-in-a-haystack to expose nerfed caching or heavy quantization.
- **ECE Leadership & Engagement Fellow** (CMU, Aug 2026 – present): Own budgeting, marketing, and outreach for ECE student programs — company-sponsored hackathons, alumni events, and tech event outings.