Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Large Language Model Systems, Machine Learning, Data Science for Software Engineering

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Large Language Models in Practice, Software Engineering, Operating Systems

WORK EXPERIENCE

NewsBreak **May – August 2026**
AI Engineer Intern *Mountain View, CA*

- Built a personalized LLM alert pipeline as a Kubernetes CronJob (embed → LLM-score → LLM dedupe) that logs why each item fired; ran it in production shadow on 4,100+ decisions for 11 days, cutting filler 34% → 4%.
- Hand-graded alerts and found non-events and real events both scored a median 89–90, so added a categorical happened / speculative / no-event gate; checked it on a 127-alert holdout and rejected a rule that caught 91% of filler but blocked the day's two highest-reach real alerts.
- Built an SSE LLM search API: deterministic RAG on the first question, a capped tool-calling agent on follow-ups (5 tools, ≤3 hops), and server-resolved citations the model cannot author; entity resolution got 20/20 logged queries right that used to return nothing, scored by an eval harness on the real catalog.
- Sole backend author for SportsBreak (∼6,000 installs): built the FastAPI / MongoDB service from its first commit to 89 REST endpoints on AWS EKS, with ∼1,450 offline tests gating every merge in CI (ruff, pyright, pytest) and a MongoDB fake that reproduces a production-cluster failure as a regression guard.
- Moved nine third-party API integrations server-side so the iOS binary ships zero API keys; set up Prometheus / Grafana and fixed a 5-day silent CronJob outage with job and scheduler deadlines.
- Diagnosed a 65% failure rate in another team's LLM pipeline, via Langfuse traces, as an unbounded reasoning loop hitting a 16,384-token ceiling; sized a 2,048-token cap against p95 1,197 tokens.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Designed database schemas and REST APIs for LangChain LLM features (resume parsing, improvement suggestions, AI job matching), and reviewed the API structure.
- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; broke down user stories, led PR reviews, and mentored interns on Git workflows and testing.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview app (React, Next.js, Supabase); worked with recruiters so AI scoring matched real hiring, and iterated with clients until three mid-sized companies signed.

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher

- Mapped natural-language instructions to formal safety specs with multi-turn LLMs and implemented Bayesian estimators for real-time road-friction adaptation (LuGre tire models).

SKILLS

Languages	Python, TypeScript, JavaScript, SQL, Bash, Go
AI & Evaluation	LLM tool calling / agents, RAG, hybrid retrieval, Langfuse, LangChain, PyTorch
Backend & Data	FastAPI, Pydantic, REST APIs, MongoDB, PostgreSQL, Redis Streams
Ops & Testing	Git, GitHub Actions, CI/CD, pytest, Docker, Linux, Kubernetes (EKS), Prometheus
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and benchmarked KV caching, long context, and needle-in-a-haystack to expose nerfed caching and quantization.
- **my-av:** Built a PyTorch CNN-GRU that predicts steering and speed from 1s video, with a data pipeline over 33 hours of Comma2k19 (shuffling, temporal windowing, frame processing).