Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Foundations of Software Engineering, Software Refactoring

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Algorithms, Data Structure & Programming III, Operating Systems

WORK EXPERIENCE

NewsBreak **May – August 2026**
Backend Engineer Intern *Mountain View, CA*

- Shipped SportsBreak to the App Store as sole backend author and primary iOS engineer; built the Python FastAPI / MongoDB service from its first commit to 89 REST endpoints and 7 Kubernetes CronJobs, with ∼1,450 offline tests gating every merge in CI (ruff, pyright, pytest).
- Stood up the containerized service on production AWS EKS (Docker, Helm, Jenkins) on day 3 with HPA autoscaling and Prometheus / Grafana monitoring; debugged a DNS → TLS → admission-policy chain that blocked iOS clients on public networks and drove a public-domain fix.
- Fixed a 5-day silent CronJob outage (stuck ImagePullBackOff + Forbid concurrency) with job and scheduler deadlines; traced a 65% failure rate in a shared Redis Streams LLM pipeline to an unbounded reasoning loop.
- Fixed an AI-agent-written cache that never hit (keyed per user, 13.2 MB each) by sharing one 10.9 MB copy as the pod neared its memory limit; cut the World feed from 25–30 client requests to one call.
- Built a real-time alert pipeline as a Kubernetes CronJob (ingest → embed → LLM-score → dedupe → push in ∼2 min); ran it in production shadow for 11 days before launch, cutting filler alerts from 34% to 4%.
- Migrated nine third-party integrations server-side so the app ships zero third-party API credentials; wrote 34 contract-first API specs and ∼10,000 lines of docs, and onboarded the team's second engineer.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to take a React, Next.js, FastAPI, and Supabase web app from 0→1 and ship it on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.
- Designed and reviewed the database schemas and REST API structure for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase, PostgreSQL); owned auth and the database and iterated with clients until three 200–500-person companies signed.

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher

- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty, including ice.

SKILLS

Languages	Java, Python, Go, C, SQL, Swift, TypeScript, JavaScript, Bash
Backend & Data	REST APIs, FastAPI, Pydantic, asyncio, MongoDB, PostgreSQL, MySQL, Redis Streams
Cloud & Testing	AWS, GCP, Kubernetes (EKS), Docker, Helm, Jenkins, pytest, Unit Testing, Git
AI Tooling	Claude Code / AI coding tools, LLM tool calling / agents, Langfuse, RAG
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **Operating Systems (UW–Madison):** Wrote a multi-threaded C quicksort with pthreads, mutexes, and condition-variable completion signaling; contributed xv6 system-call entry points and a `clone` wrapper.
- **Wisconsin Autonomous:** Integrated wheel-speed encoders on the competition rover and streamed real-time data to the onboard NVIDIA Jetson.
- **my-av:** Built a PyTorch CNN-GRU that predicts steering and speed from 1 s video; trained on GCP Vertex AI.