Lordphone Wen

Mountain View, CA

lordphonewen@gmail.com · (608) 217-8920 · linkedin.com/in/lordphone-wen · github.com/lordphone

EDUCATION

Carnegie Mellon University **Expected May 2027**
M.S. in Software Engineering | GPA: 4.0 / 4.0
Selected Coursework: Distributed Systems, Large Language Model Systems, Machine Learning

University of Wisconsin–Madison **May 2024**
B.S. in Computer Science & B.S. in Computer Engineering | GPA: 3.72 / 4.0
Selected Coursework: Operating Systems, Computer Architecture, Algorithms

WORK EXPERIENCE

NewsBreak **May – August 2026**
Engineering Intern *Mountain View, CA*

- Stood up the containerized Python service on production AWS EKS (Docker, Helm, Jenkins) on day 3; traced a DNS → TLS → admission-policy chain blocking iOS clients on public networks to a public-domain fix.
- Set up Prometheus metrics and Grafana dashboards for the backend's EKS nodes; configured HPA autoscaling (2–4 pods at 70% CPU) and a two-stage amd64 Docker image after a C extension crashed on arm64.
- Added a ruff / pyright / pytest GitHub Actions gate after tests had sat red on `main` (Jenkins only built the image), and caught a stale ArgoCD app with prune and self-heal on that would have wiped Jenkins deploys.
- Fixed a 5-day silent Kubernetes CronJob outage (non-terminal ImagePullBackOff + Forbid concurrency) with paired job and scheduler deadlines; traced a 65% failure rate in another team's LLM pipeline via Langfuse to an unbounded reasoning loop and a Helm env value shadowing the code default.
- Fixed a cache that never hit (keyed per user, 13.2 MB each) by sharing one 10.9 MB copy as the 512 MiB pod neared its memory limit; cut the World feed from 25–30 client requests to one backend call.
- Sole backend author for SportsBreak (∼6,000 installs): built the FastAPI / MongoDB service from its first commit to 89 REST endpoints and 7 CronJobs, including an LLM alert pipeline (∼1,200 LLM calls/day) run in production shadow for 11 days before launch.

Haddee Education **June – September 2025**
Engineering Lead *San Diego, CA*

- Led an 8-intern team through Agile sprints to ship a React, Next.js, FastAPI, and Supabase web app on GCP App Engine; led PR reviews and mentored interns on Git workflows and testing.
- Designed database schemas and REST APIs for LangChain LLM features: resume parsing, improvement suggestions, and AI job matching.

ZenAI — MiraclePlus (Formerly YC China) F23 **September – December 2023**
Full Stack Developer *Beijing, China*

- Sole engineer on a VC-backed startup's LLM mock-interview web app (React, Next.js, Supabase, PostgreSQL); owned auth and the database and iterated with clients until three 200–500-person companies signed.

Control & Learning Group — Carnegie Mellon University **February – June 2026**
Graduate Researcher

- Developed an Online Adaptive Probabilistic Safety Certificate (PSC) with PhD researchers; outperformed Adaptive MPC in 3-DOF vehicle simulations under environmental uncertainty.

SKILLS

Cloud & Infra	AWS, GCP, Kubernetes (EKS), kubectl, Helm, Docker, Terraform, Linux, Bash
CI/CD & Monitoring	GitHub Actions, Jenkins, ArgoCD, CI/CD, Prometheus, Grafana, Langfuse
Languages	Python, Go, C, SQL, Java, TypeScript
Backend & AI	FastAPI, REST APIs, MongoDB, PostgreSQL, RAG, LLM tool calling / agents
Certifications	AWS Certified Solutions Architect – Associate, AWS Certified Cloud Practitioner

PROJECTS

- **codingplans:** Building a site that rates LLM coding plans; bought Z.AI, MiniMax, and Kimi plans and benchmarked KV caching, long context, and needle-in-a-haystack to expose nerfed caching and quantization.
- **Operating Systems (UW–Madison):** Wrote a multi-threaded C quicksort with pthreads and mutexes; contributed xv6 system-call entry points.
- **Wisconsin Autonomous:** Streamed real-time rover wheel-speed encoder data to the onboard NVIDIA Jetson.